

Information Security and Privacy

David Jelenc

Administrative info

- 3 lab session slots
 - Check official Timetable
- Office hours
 - Slot: TBD and published on Moodle
 - Office: R3.50, third floor, first bridge on the left
 - Welcome at any time, but email first
- Point of contact
 - Forum in Moodle (preferred)
 - david.jelenc@fri.uni-lj.si

Syllabus (1/2)

- Using cryptoprimitives to develop secure applications using cryptographic libraries (JACL)
- Topics (by weeks)
 - Secrecy (stream and block ciphers)
 - Integrity (Message digests, MACs) & authenticated encryption
 - Public-key encryption, steganography,
 - Key exchange protocols, digital signatures, key-derivation
- Prerequisite: Java programming (basics)
- At the end: Midterm 1

Syllabus (2/2)

- Using various tools to secure and securely administrate computer systems (Ubuntu Linux)
- Topics (by weeks)
 - Firewalls with Netfilter/IPtables
 - Firewalls continued, NAT, routing
 - Secure shell with OpenSSH
 - Virtual private networking with IPsec
 - Authentication, authorization and accounting with FreeRADIUS
- Prerequisite: Computer networking & Linux basics
- At the end: Midterm 2

Grading

- **Weekly homework assignments**
 - Needs to be handed in by the end of the week
 - Pass/Fail grading: need at least 7/9 to pass
- **Homework challenges**
 - Optional homework assignments; allows students to receive extra credit
- **Midterms**
 - During the semester, one after the completion of each syllabus topic
 - Format: quiz + programming assignment
 - Quiz [50%]: 10 multiple-choice/closed-form-type questions
 - Programming assignment [50%]:
 - A short seminar-like task, similar to those at lab sessions
 - The solution has to be *defended*; the grading is done in the presence of the student
 - Programming is done on classroom computers, open-book style, Internet disconnected
- **Final lab total**
 - $\text{Points} = \text{MAX}(0.5 * \text{MidTerm1} + 0.5 * \text{MidTerm2} + \text{ExtraPoints}, 100)$

Tentative grading dates

- **Dates are tentative, may still change**
- **Midterm 1**
 - Week Nov 25 to Dec 1
- **Midterm 2**
 - Week Jan 6 to Jan 12

Communication Secrecy

Contents


- Introduction
- Stream ciphers
 - Perfect secrecy
 - One time pad (OTP)
 - Pseudorandom generators (PRG)
 - Semantic security for one-time keys
- Block ciphers
 - Pseudorandom functions and permutations (PRFs, PRPs)
 - Modes of Operation
- Semantic security for many-time keys
- Summary

Introduction: providing confidentiality

- We'd like to provide confidential communication
 - Only the intended recipient(s) should be able to read the data



- Two types of encryption and decryption
 - Symmetric ciphers
 - Asymmetric ciphers



Symmetric Ciphers

- A cipher defined over (K, M, C) is a pair of “comp. eff.” algorithms (\mathbf{E}, \mathbf{D}) , where

$$\mathbf{E}: K \times M \rightarrow C$$

$$\mathbf{D}: K \times C \rightarrow M$$

s. t. for all k in K and m in M :

$$\mathbf{D}(k, \mathbf{E}(k, m)) = m$$

- \mathbf{E} is often randomized, \mathbf{D} is always deterministic

Perfect Secrecy

- What is a “secure” cipher?
 - Shannon: Cipher text should reveal “no information” about the plain text
- A cipher (E, D) over (K, M, C) has **perfect secrecy** if for all $m_0, m_1 \in M$ ($|m_0|=|m_1|$) and for all $c \in C$
$$\Pr [E(k, m_0) = c] = \Pr [E(k, m_1) = c]$$
where $k \in K$ is randomly chosen
 - Given cipher text c , one cannot tell whether c is a cryptogram of m_0 or m_1

One Time Pad

- Vernam (1917)
 - $M = C = K = \{0, 1\}^n$
 - $E(k, m) = k \oplus m$
 - $D(k, c) = k \oplus c$
- Features
 - Given a truly random key, OTP has ***perfect secrecy***
 - Key has to be ***random*** and it must be used ***only once***
 - Impractical: Shannon shows that perfect secrecy requires keys to be at least as long as the plain text

Pseudo Random Generator

- Idea: Replace a “random” with a “pseudorandom” key
$$\mathbf{G}: \{0, 1\}^s \rightarrow \{0, 1\}^n \quad \text{where } n \gg s$$
- Pseudo Random Generator (PRG) is a function \mathbf{G} that maps *seed space* to *key space*
 - Is “efficiently” computable by a deterministic algorithm
 - Its output (keys) “looks random”
- Stream ciphers
 - $\mathbf{E}(k, m) := m \oplus \mathbf{G}(k)$
 - $\mathbf{D}(k, c) := c \oplus \mathbf{G}(k)$
- Examples: RC4, CSS, eStream, Salsa 20
- Can stream ciphers have perfect secrecy, why?

Stream Ciphers: perfect secrecy?

- Stream ciphers cannot have perfect secrecy
 - Keys (seeds) are shorter than messages
- Can stream ciphers ever be secure?
 - Need a new definition of security
 - Security will depend on PRG used

Pseudo Random Generators: defs

- **Statistical test** is an algorithm $A: \{0, 1\}^n \rightarrow \{0, 1\}$
 - Returns 1 if it *thinks* the input string is random, 0 otherwise
- **Advantage** of st. test A against PRG G :
$$\text{Adv}_{\text{PRG}}[A, G] = \left| \Pr_{k \xleftarrow{R} K} [A(G(k)) = 1] - \Pr_{r \xleftarrow{R} \{0, 1\}^n} [A(r) = 1] \right|$$
 - If *close* to 0, A cannot distinguish G from random
 - Otherwise, A can distinguish G from random
- **Def.** A PRG G is secure, if for all eff. stat. tests A :
$$\text{Adv}_{\text{PRG}}[A, G] \text{ is negligible.}$$

Negligible? Assume less than 2^{-80}

Pseudo Random Generators: defs

- Def: A PRG is **unpredictable** if given an initial sequence of bits (a prefix), one cannot *efficiently* predict the next bit (with probability higher than $\frac{1}{2} + \epsilon$)
- Thm: A PRG is secure iff. it is unpredictable.
- In practice
 - Unknown if there are provably secure PRG
 - But we have heuristic candidates

Perfect secrecy, threat model

- (Recall) A cipher (E, D) over (K, M, C) has **perfect secrecy** if for all $m_0, m_1 \in M$ ($|m_0|=|m_1|$) and for all $c \in C$
$$\Pr [E(k, m_0) = c] = \Pr [E(k, m_1) = c]$$
where $k \in K$ is randomly chosen
 - Given cipher text c , one cannot tell whether c is a cryptogram of m_0 or m_1
- **Threat model:** basis for reasoning about security
 - **Adversary's power:** what can she do
 - **Adversary's goal:** what is she trying to achieve

Semantic security: def


(for one-time key; adv. sees only one CT)

- Adversary's power: **observe one ciphertext**
 - Every message is encrypted with its own key; a particular key is used only once
- Adversary's goal: **learn about the plaintext**

Semantic security: def

(for one-time key; adv. sees only one CT)

- For $b \in \{0, 1\}$ define experiments $\text{EXP}(b)$ as



- Def: $\zeta = (E, D)$ is **semantically secure** if for all eff. adversaries A $\text{Adv}_{\text{ss}}[A, \zeta]$ is negligible.

$$\text{Adv}_{\text{ss}}[A, \zeta] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$

Semantic security

- Informally
 - A cipher has **semantic security** if given only cipher text, an attacker cannot *practically* derive any information about the plain text
- **Thm:** Given a secure PRG, derived stream cipher is semantically secure

Final thoughts

- Two-time pad attack
 - Never use stream-cipher key to encrypt more than one message
 - later we show a secure a multi-message exchange

$$c_1 \leftarrow m_1 \oplus \mathbf{G}(k)$$


$$c_2 \leftarrow m_2 \oplus \mathbf{G}(k)$$

$$m_1 \oplus m_2 \leftarrow c_1 \oplus c_2$$

- Redundancy in natural languages and in encoding schemes (ASCII, UTF-8, ...) to separate $m_1 \oplus m_2 \rightarrow m_1, m_2$
- <http://www.crypto-it.net/eng/attacks/two-time-pad.html>

Final thoughts

- Two-time pad attack



Final thoughts

- Malleability
 - Modifications to CT are not detected and have predictable impact on the plain text

Encrypt: $c \leftarrow m \oplus k$


Modify: $c' \leftarrow c \oplus \mathbf{p}$

Decrypt: $m' \leftarrow c' \oplus k$

- What is the relation between m and m' ?


Block Ciphers

- Notable examples
 - 3DES: $n = 64$ bits, $k = 168$ bits
 - AES: $n = 128$ bits, $k = 128, 192, 256$ bits



Block Ciphers: Built by iteration

- $R(k, m)$ is a round function
 - 3DES ($n = 48$)
 - AES ($n = 10$)



Abstracting BC: PRF and PRP


- Pseudo Random Function (PRF) defined over (K, X, Y) :
 $F: K \times X \rightarrow Y$
 - We can evaluate $F(k, x)$ efficiently
- Pseudo Random Permutation (PRP) defined over (K, X) :
 $E: K \times X \rightarrow X$
 - We can evaluate $E(k, x)$ efficiently
 - $E(k, -)$ has an inverse
 - We have an efficient inversion algorithm $D(k, x)$
 - (All PRPs are PRFs.)

Secure PRF

- Let $F : K \times X \rightarrow Y$ be a PRF
 - $Funs[X, Y]$ the set of all functions from X to Y
 - $S_F = \{F(k, -) : \forall k \in K\} \subseteq Funs[X, Y]$


Intuitively

- A PRF is secure if a random function in $Funs[X, Y]$ is indistinguishable from a random function in S_F
- Believed to be secure PRPs:
 - AES, 3DES, Blowfish



Secure PRF (def.)

- For $b \in \{0, 1\}$ define experiment $\text{EXP}(b)$ as




- Def: F is a secure PRF if for all eff. adversaries A $\text{Adv}_{\text{PRF}}[A, F]$ is negligible.

$$\text{Adv}_{\text{PRF}}[A, F] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$

Secure PRP (def.)

- For $b \in \{0, 1\}$ define experiment $\text{EXP}(b)$ as



- Def: E is a secure PRP if for all eff. adversaries A $\text{Adv}_{\text{PRP}}[A, E]$ is negligible.

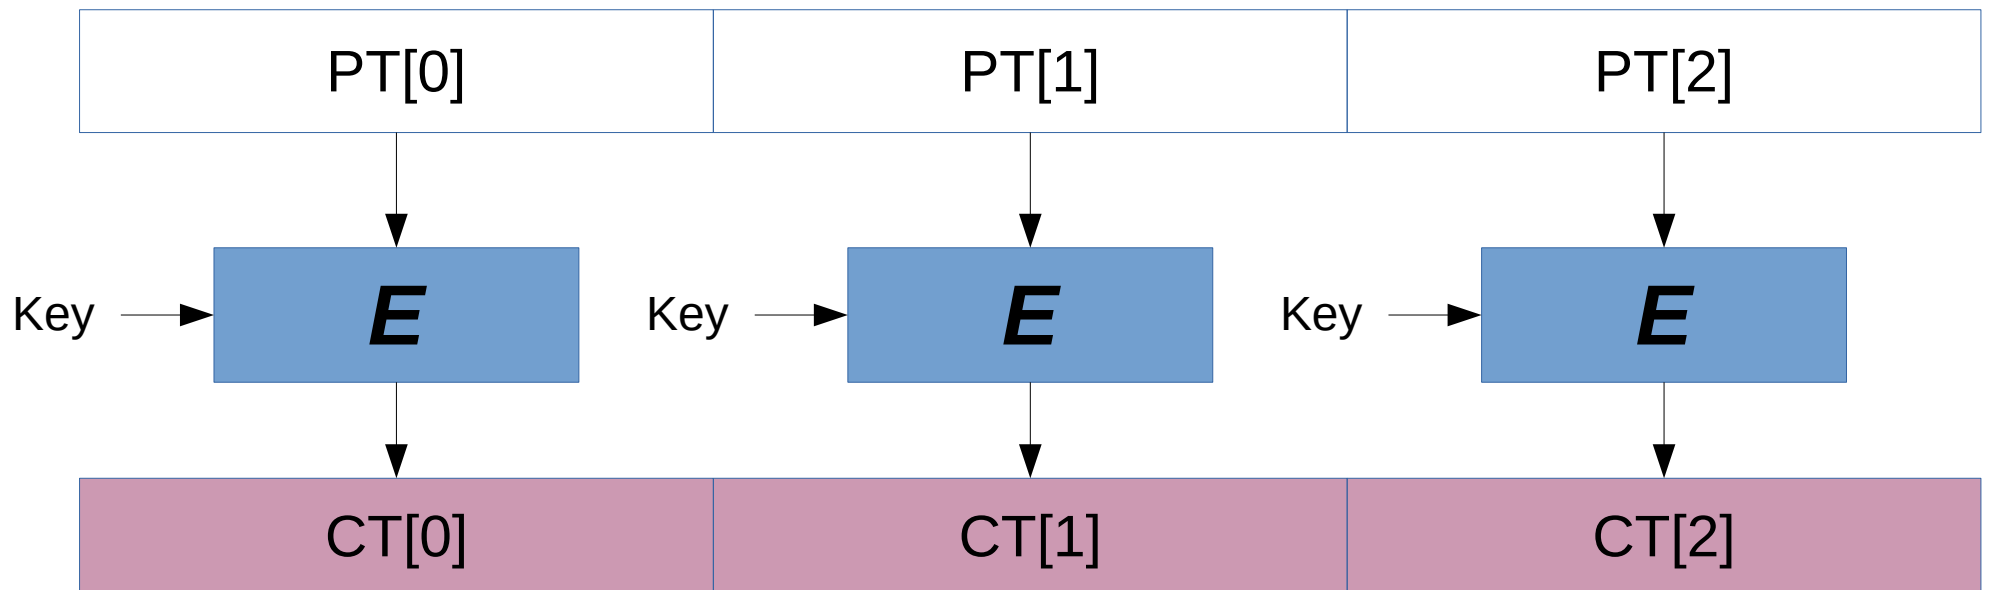
$$\text{Adv}_{\text{PRP}}[A, E] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$

Block Ciphers: Modes of Operation

- Goal: How do we build a secure encryption from secure PRP (e.g. AES)
 - A PRP encrypts a single data block. How do we encrypt larger data?
- Semantic security (still for one-time key only)
 - Adversary's power: **observe one ciphertext**
 - Adversary's goal: **learn about plaintext**



MO: Electronic Code Book

- “Solution” Electronic Code Book (ECB):
 - Split the data into blocks
 - if needed, extend the last block with padding bits
 - Independently encrypt each block




MO: Electronic Code Book

- Problem: If $PT[0] == PT[1]$, then $CT[0] == CT[1]$
 - If two plaintext blocks are the same, so are the corresponding ciphertexts blocks



How does the Adversary win the semantic security game against ECB?

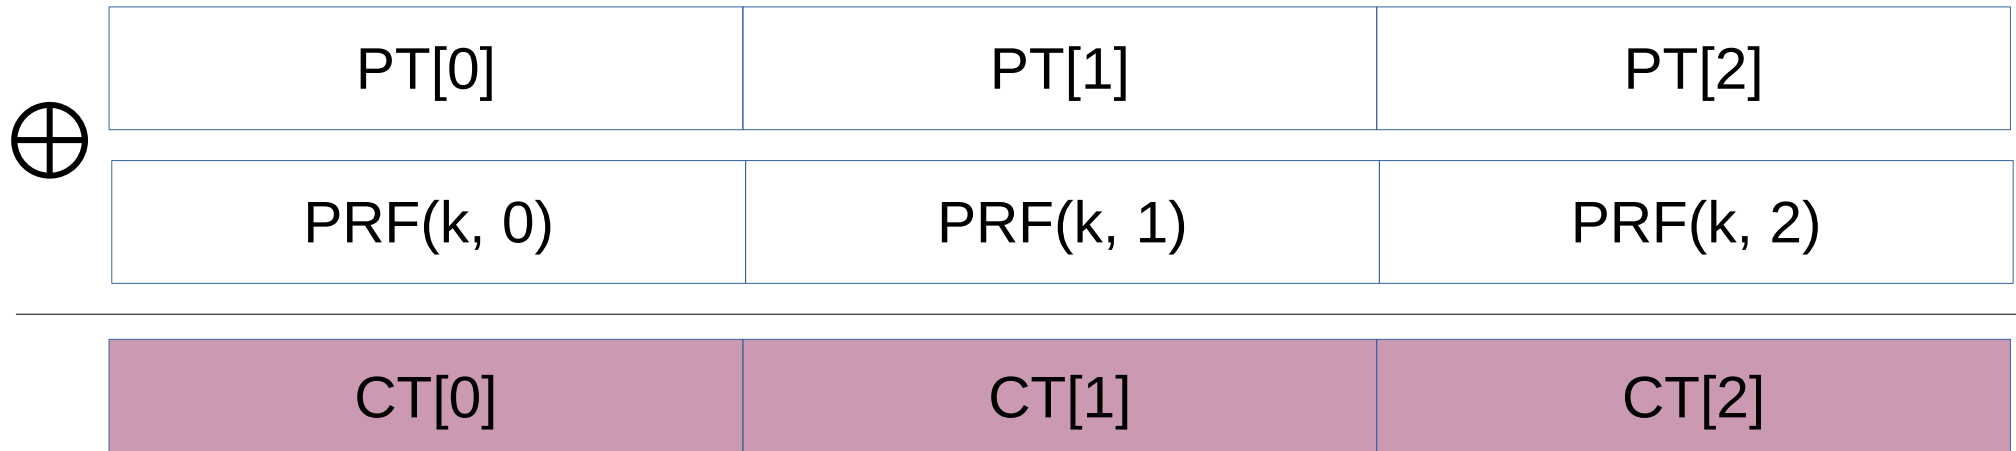
ECB is not semantically secure



- What should be m_0, m_1 ?
- What should the adversary output?

MO: Deterministic counter mode

- Deterministic counter from a pseudorandom function (PRF)



- Creates a stream cipher from a PRF
- Secure (but only for encrypting a single message which may consists of multiple blocks)


Semantic security for many-time key

- Key is used more than once: adversary sees many CTs encrypted with the same key
- Adversary's power: **chosen-PT attack (CPA)**
 - Can obtain the encryption of any message of her choice
- Adversary's goal: **break semantic security**
 - Learn about the PT from the CT

Semantic security for CPA (def)

(for many-time key)


- Let $\zeta = (E, D)$ be a cipher defined over (K, M, C)
- For $b \in \{0, 1\}$ define experiments $\text{EXP}(b)$ as



- Def: $\zeta = (E, D)$ is **semantically secure under CPA** if for all eff. adversaries A $\text{Adv}_{\text{CPA}}[A, \zeta]$ is negligible.
- $$\text{Adv}_{\text{CPA}}[A, E] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$

Ciphers insecure under CPA

- Suppose a cipher is deterministic
 - Given some message m , the cipher always produces the same ciphertext




- Implication
 - An attacker can learn that two encrypted elements (files, packets, ...) are the same

Ciphers insecure under CPA

- If a key is to be used multiple times, the encryption should be **non-deterministic**:
 - Encrypting the same PT twice, must produce different CTs
- Solutions
 - Randomized encryption
 - Nonce-based encryption

Non-deterministic encryption

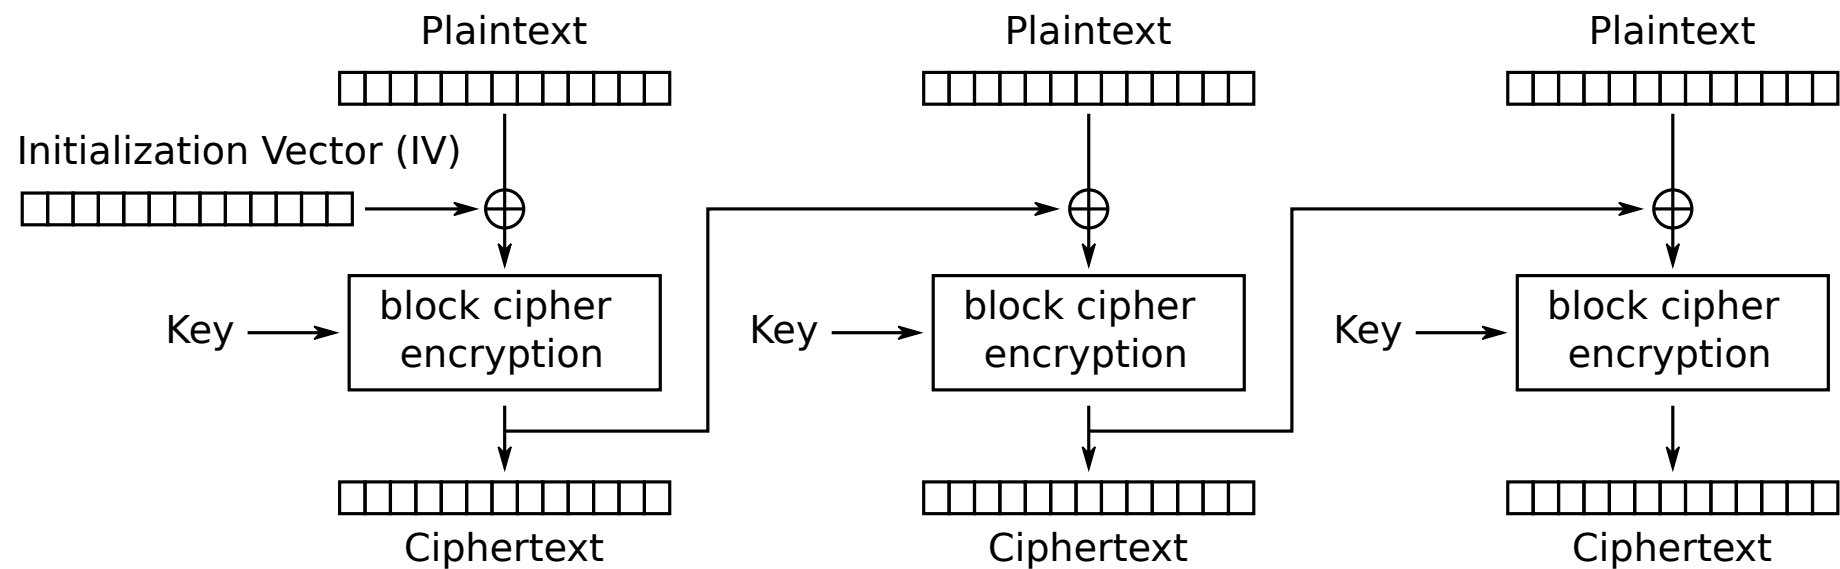
- Nonce ***n***: a value that changes from message to message
 - Pair (*key*, *n*) must never repeat
- Method 1: Nonce is a random value (AES-CBC)
- Method 2: Nonce is a counter (AES-CTR)




- CPA system should be secure even when the adversary chooses nonces

Modes of Operation: CBC

- Randomize the encryption with an initialization vector (IV)
 - Sent unencrypted
 - Must generate new random IV for every message: pair **(key, IV)** must never repeat
 - IV must be unpredictable
- Forces encryption to be sequential
 - Decryption may be parallelized




Cipher Block Chaining (CBC) mode encryption




Cipher Block Chaining (CBC) mode decryption

Counter Mode

- The random element is a counter
 - Or a combination of a random IV and a counter
 - The combination must not repeat for the lifetime of the key
- Encryption and decryption can be done in parallel
- In effect, creates a stream cipher out of a block cipher



Counter (CTR) mode encryption



Counter (CTR) mode decryption

Summary

- Two security notions
 - Semantic security against one-time CPA
 - Semantic security against many-time CPA
- Only covered secrecy against passive attackers
 - Adversaries can see, but not modify cipher text
 - We'll cover integrity next week

Goal \ Power	Power	
	One-time key	Many-time key (CPA)
Semantic security	Stream-ciphers Deterministic CTR-mode	Rand CBC Rand CTR-mode

Integrity

Contents

- Introduction
- MAC Definition
 - PRF
 - Secure PRF \rightarrow Secure MAC
- ECBC-MAC
- Cryptographic hash functions
 - Collision resistance
 - MACs from CR
 - Merkle-Damgard iterative construction
- HMAC

Introduction

- Integrity: maintaining accuracy and completeness of data
- Goal
 - Prevent adversary from modifying data
 - More feasible: detect if data has been altered
- Examples
 - Protecting files on disks
 - Assuring installation of correct software
 - Assuring the delivered packet has not been tempered with in traffic

Message Authentication Code



$MAC I = (S, V)$ defined over (K, M, T) is a pair of algs.:

$$S : K \times M \rightarrow T$$

$$V : K \times M \times T \rightarrow \{0, 1\}$$

$$|M| \gg |T|$$

such that

$$\forall k \in K, m \in M: V(k, m, S(k, m)) = 1$$

Is a shared secret required?

- Is all these secrecy required?
- Could we not just simply use
 - MD-5 or
 - SHA-{1,2,3} or
 - CRC?

Secure MAC


- Attacker's power: **Chosen message attack**
 - For $m_1 \dots m_q$ attacker is given $t_i = S(k, m_i)$
- Attacker's goal: **Existential forgery**
 - Produce a **new** valid (m, t) s. t.

$$(m, t) \notin \{(m_1, t_1) \dots (m_q, t_q)\}$$

Implications

- attacker cannot produce a valid tag for a new message
- given (m, t) attacker cannot produce (m, t') for $t \neq t'$

Secure MAC (def)



$b \in \{0, 1\}$

$b=1$ if $V(k, m, t)=1$ and $(m, t) \notin \{(m_1, t_1) \dots (m_q, t_q)\}$

$b=0$ otherwise

$I = (S, V)$ is a **secure MAC** if for all “efficient” adversaries A

$$\text{Adv}_{\text{MAC}}[A, I] = \Pr[\text{Chal. outputs } 1]$$

is “negligible”.

Secure MAC


- Negligible?
 - Assume less than 2^{-80}
- Suppose a $S(k, m)$ computes 10-bit tags
 - Is such a MAC secure, why?

(Recall) Secure PRF

- Let $F : K \times X \rightarrow Y$ be a PRF
 - $Funs[X, Y]$ the set of all functions from X to Y
 - $S_F = \{F(k, -) : \forall k \in K\} \subseteq Funs[X, Y]$


Intuitively

- A PRF is secure if a random function in $Funs[X, Y]$ is indistinguishable from a random function in S_F



(Recall) Secure PRF (def.)

- For $b \in \{0, 1\}$ define experiment $\text{EXP}(b)$ as



- Def: F is a secure PRF if for all eff. adversaries A $\text{Adv}_{\text{PRF}}[A, F]$ is negligible.

$$\text{Adv}_{\text{PRF}}[A, F] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$

Secure PRF \rightarrow Secure MAC

- For a PRF $F : K \times X \rightarrow Y$ define MAC $I_F = (S, V)$

$$S(k, m) := F(k, m)$$

$$V(k, m, t) := \begin{cases} 1 & t = F(k, m) \\ 0 & \text{otherwise} \end{cases}$$

- **Thm.** If F is a secure PRF and $1/|Y|$ is negligible (i.e. $|Y|$ is sufficiently large), then I_F is a secure MAC.


Truncating MACs based on PRFs

- Lemma: Suppose $F: K \times X \rightarrow \{0, 1\}^n$ is a secure PRF. So is $F_t(k, m) := F(k, m)[1 \dots t]$ for all $1 \leq t \leq n$
- If (S, V) is a MAC based on a secure PRF that outputs n -bit tags, then the truncated MAC that outputs w bits is also secure.
 - As long as 2^{-w} is still negligible

Examples of secure MAC

- AES (or any secure PRF)
 - A secure MAC for 16-byte (128-bit) messages
- Longer messages?
 - **CBC-MAC**
 - **HMAC**
- Both convert a small-PRF into a big-PRF

ECBC-MAC



Hash-MAC (HMAC)

- Built from *collision resistance*
- Let $H : M \rightarrow T$ be a hash function $|M| \gg |T|$
- A **collision** for H is a pair $m_0, m_1 \in M$ such that:
$$H(m_0) = H(m_1) \text{ and } m_0 \neq m_1$$
- Function H is **collision resistant** if for all *explicit* “eff.” algs. A $\text{Adv}_{\text{CR}}[A, H]$ is negligible.

$$\text{Adv}_{\text{CR}}[A, H] := \Pr[A \text{ outputs collision for } H]$$

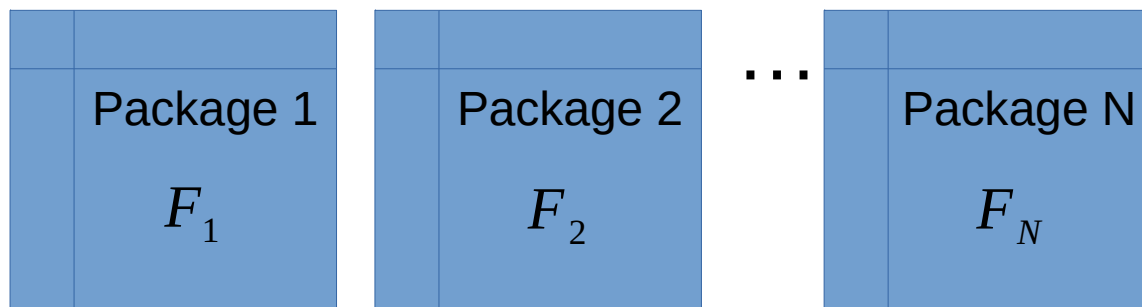
- Example: SHA-256

MAC from CR

- Let $I = (S, V)$ be a MAC for short messages over (K, M, T) (e.g. AES)
- Let $H : M^{\text{BIG}} \rightarrow M$
- Def: $I^{\text{BIG}} = (S^{\text{BIG}}, V^{\text{BIG}})$ over (K, M^{BIG}, T) as:
$$S^{\text{BIG}}(k, m) := S(k, H(m))$$
$$V^{\text{BIG}}(k, m, t) := V(k, H(m), t)$$
- **Thm.** If I is a secure MAC and H is collision resistant, then I^{BIG} is a secure MAC.
- Example: $S(k, m) := \text{AES}_{\text{2-block-CBC}}(k, \text{SHA-256}(m))$

Example: Integrity using CR hash

- Protecting software packages (Linux distros)



**READ-ONLY
public space**

$H(F_1)$
 $H(F_2)$
 $H(F_N)$

- User downloads a package and verifies it using hashes in public space
 - If H is collision resistant, the attacker cannot modify packages without being detected
- We require no shared secret, but we need a read-only public space

Generic attack on CR

- Let $H : M \rightarrow \{0,1\}^n$ be a hash function $|M| \gg 2^n$
- Generic algorithm to find a collision
 - 1) Chose $\sqrt{2^n} = 2^{\frac{n}{2}}$ random messages: $m_1 \dots m_{2^{n/2}} \in M$ distinct w.h.p.
 - 2) For $i = 1 \dots 2^{n/2}$: compute $t_i = H(m_i)$
 - 3) Look for a collision $(t_i = t_j)$. If not found, go to 1.
- How many iterations before we find a collision?

The birthday paradox

- **Thm.** Let $r_1 \dots r_n \in [1 \dots B]$ be independent and identically distributed integers. If we sample $n = 1.2 \times \sqrt{B}$ samples from interval $[1 \dots B]$ then the probability of finding a collision is

$$\Pr [\exists i \neq j : r_i = r_j] \geq 0.5$$


- Approximation of collision probability given n samples with Taylor series

$$p(n) \approx 1 - e^{\frac{-n(n-1)}{2B}}$$

Collision
probability

Collision probabilities

$$B=10^6$$



Generic attack on CR

- Let $H : M \rightarrow \{0,1\}^n$ be a hash function $|M| \gg 2^n$
- Generic algorithm to find a collision
 - 1) Chose $\sqrt{2^n} = 2^{\frac{n}{2}}$ random messages: $m_1 \dots m_{2^{n/2}} \in M$ distinct w.h.p.
 - 2) For $i = 1 \dots 2^{n/2}$: compute $t_i = H(m_i)$
 - 3) Look for a collision ($t_i = t_j$). If not found, go to 1.
- How many iterations before we find a collision?
 - ~ 2
 - Running time $O(2^{\frac{n}{2}})$

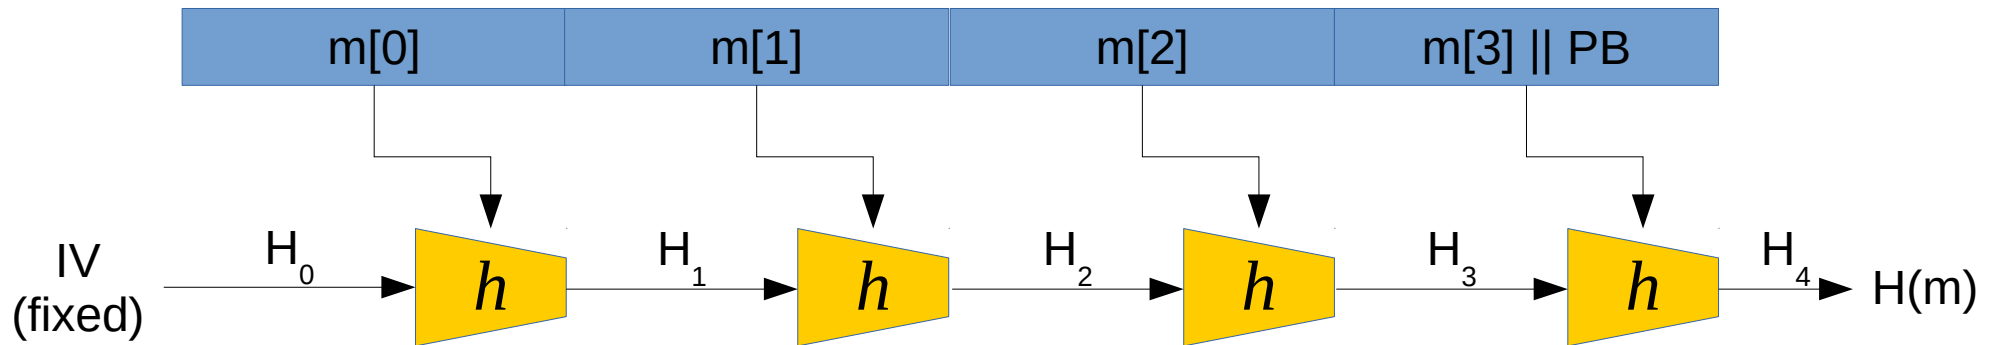
Example CR hash functions

Function	Digest (tag) size [bits]	Generic attack time
MD-5	128	2^{64}
SHA-1*	160	2^{80}
SHA-256	256	2^{128}
SHA-512	512	2^{256}
Whirpool	512	2^{256}

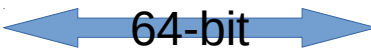
* Found collision by performing $2^{63.1}$ evaluations <https://shattered.it>

Merkle-Damgard construction

- Goal: given CR function for **short** messages, construct CR function for **long** messages




- CR for short messages (compression function) $h: T \times X \rightarrow T$
- CR for long messages $H: X^{\leq L} \rightarrow T$
- PB: padding block $10..0 \parallel \text{msg len (in bits)}$



 - If no space for PB, add an extra block
- **Thm.** If h is CR, so is H .

Compression functions

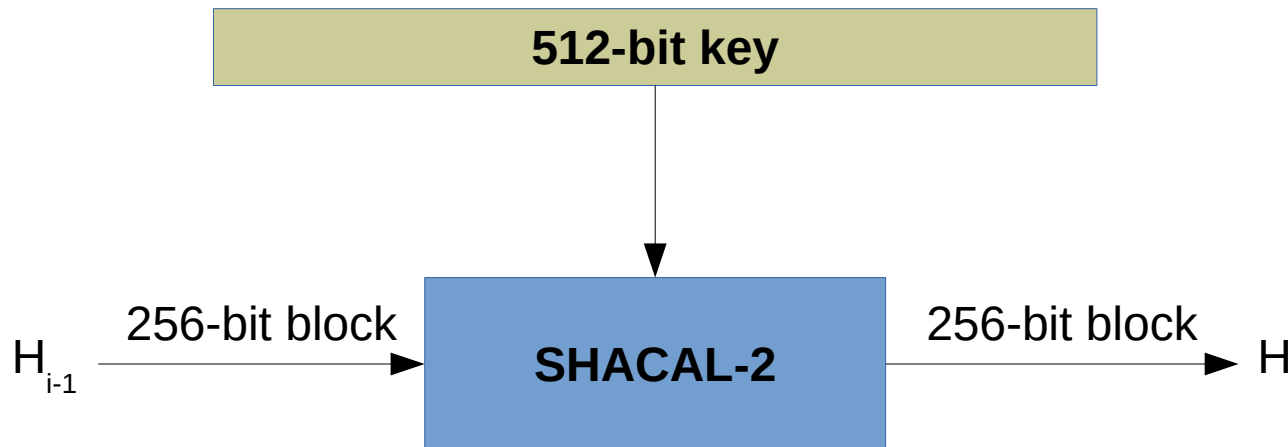
- Built from block ciphers $E : K \times \{0, 1\}^n \rightarrow \{0, 1\}^n$
- Several constructions
 - **Davies-Meyer**
$$h(H, m) := E(m, H) \oplus H$$
 - Matyas–Meyer–Oseas
 - Miyaguchi–Preneel



https://en.wikipedia.org/wiki/One-way_compression_function

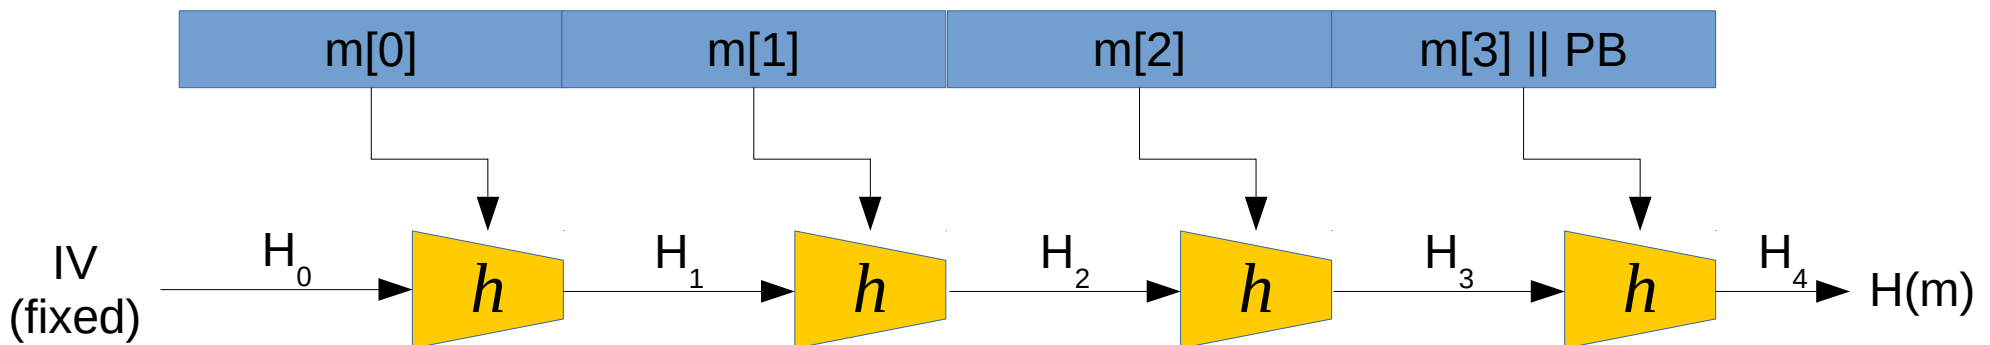
Example: SHA-256

- Merkle-Damgard iterative construction
- Davies-Meyer compression function
 - Block cipher: SHACAL-2



MAC from M-D hash func.

- Can we construct a MAC directly from H ? (e.g SHA-256)
- Naive attempt $S(k, m) := H(k \parallel m)$
 - Is it secure?



- If you knew $H(k \parallel m)$ could you compute $H(k \parallel m \parallel \text{PB} \parallel w)$ for any w ? How?
- Length-extension attack

Standardized solution: HMAC

- Most commonly used on the Internet
 - <https://tools.ietf.org/html/rfc2104>
- Given CR hash function H , define a MAC as
$$S(k, m) := H(k \oplus \text{opad} \parallel H(k \oplus \text{ipad} \parallel m))$$
 - Built from a black-box implementation of SHA-256
 - Assumed to be a secure PRF
 - TLS 1.2 requires support of HMAC-SHA1-96 (TLS 1.3 does not)

Authenticated Encryption

Contents

- Ciphertext integrity
- AE definitions
- Chosen Ciphertext Attack
- Constructions
 - Encrypt-then-MAC
 - Encrypt-and-MAC
 - MAC-then-Encrypt

Authenticated Encryption (AE)

- Everything demonstrated so far provides
 - either integrity
 - or confidentiality (security against eavesdropping)
- CPA security does not provide secrecy against active attacks (where an attacker can tamper with ciphertext)
 - ➔ If you require integrity → **MAC**
 - ➔ If you require integrity and confidentiality → **AE**

AE: Desired properties

- An authenticated encryption system $\zeta = (E, D)$ is a cipher where

as usual $E : K \times M \times N \rightarrow C$

but $D : K \times C \times N \rightarrow M \cup \{\perp\}$ $\perp \notin M$


Nonce

CT is invalid
(rejected)

- Security: the system must provide
 - **semantic security under CPA**, and
 - **ciphertext integrity**
 - an adversary cannot create a new valid CT (such that would decrypt properly)

Ciphertext integrity (def)

Let $\zeta = (E, D)$ be a cipher with message space M



$\downarrow b \in \{0, 1\}$

$b = 1$ if $D(k, c) \neq \perp$ and $c \notin \{c_1 \dots c_q\}$

$b = 0$ otherwise

Def: $\zeta = (E, D)$ has **ciphertext integrity** if for all “efficient” adversaries A : $\text{Adv}_{\text{CI}}[A, \zeta]$ is “negligible”.

$$\text{Adv}_{\text{CI}}[A, \zeta] = \Pr[\text{Chal. outputs } 1]$$

Authenticated Encryption

- Def: A cipher $\zeta = (E, D)$ **provides authenticated encryption (AE)** if it is
 - 1) semantically secure under CPA, and
 - 2) has ciphertext integrity.
- Do the following ciphers provide AE:
 - AES-CBC,
 - AES-CTR,
 - RC4?
- Why?

Authenticated Encryption

- Implication 1: Authenticity



- An attacker cannot create a new valid $c \notin \{c_1 \dots c_q\}$
- If message decrypts properly ($D(k, c) \neq \perp$), it must have come from someone who knows secret key k
 - But it could be a replay

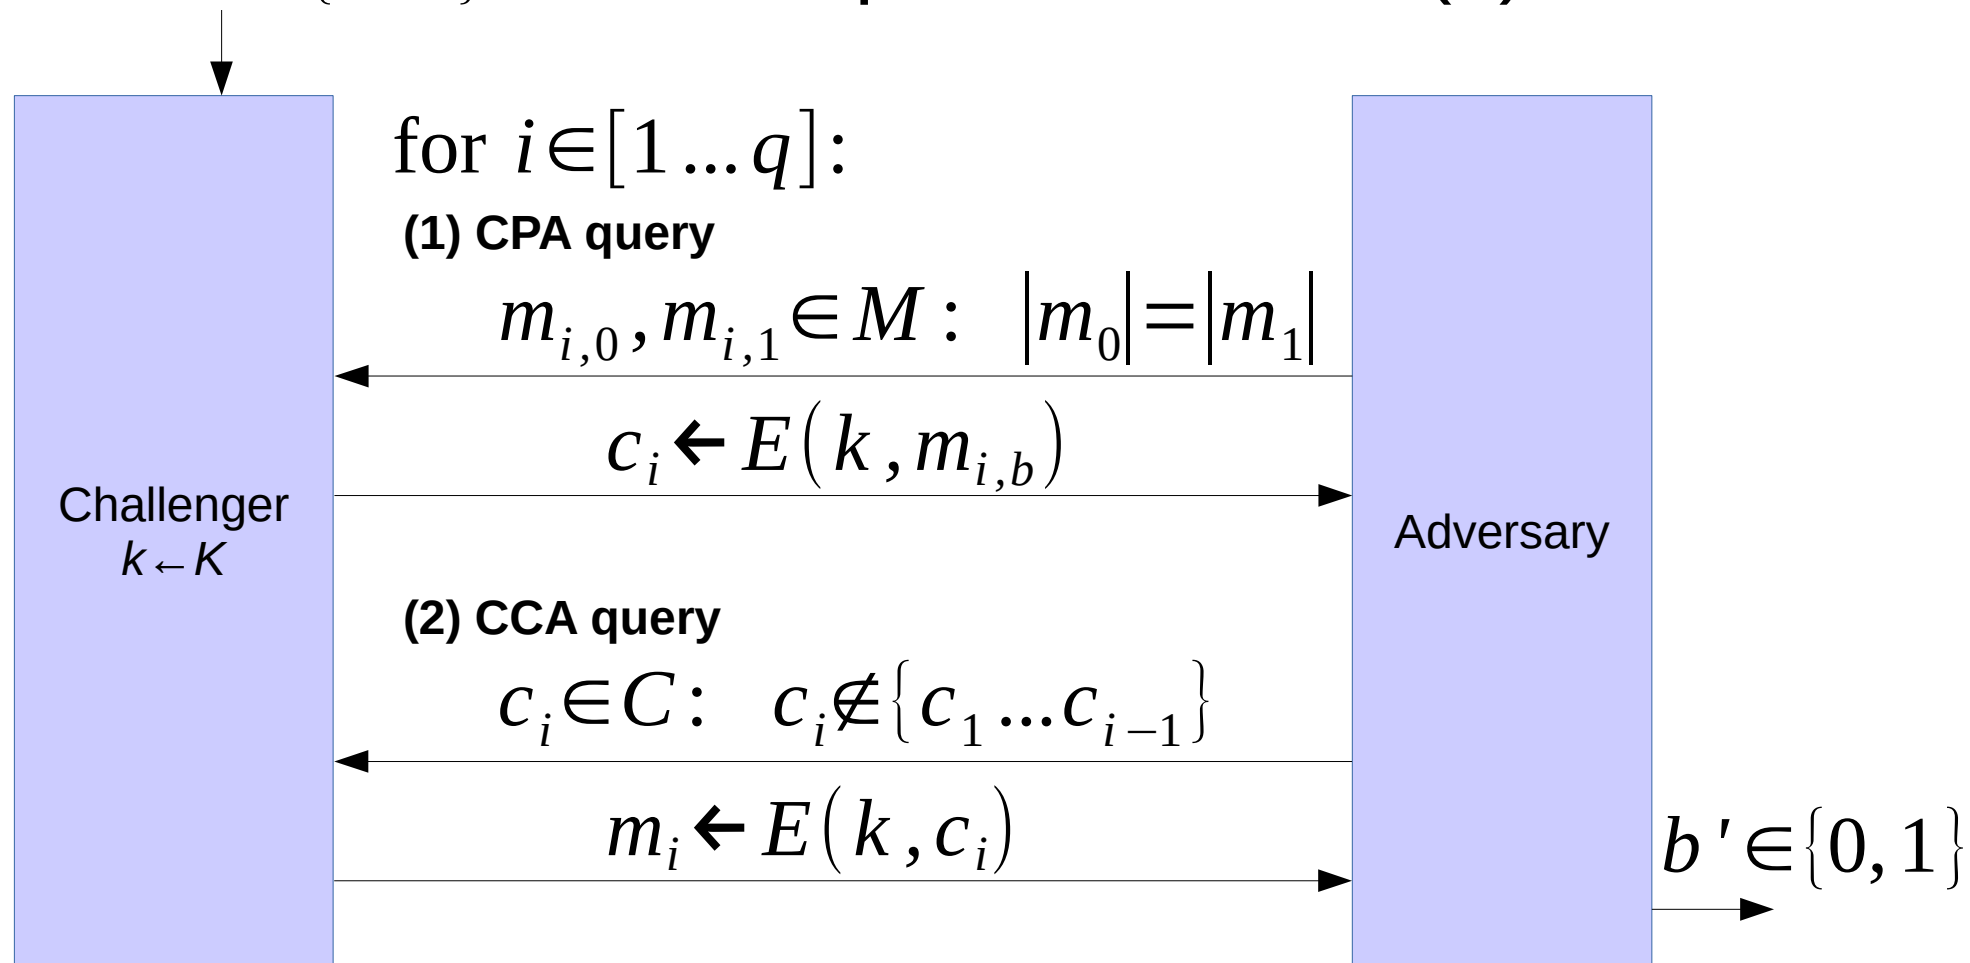
- Implication 2: Security against **chosen ciphertext attack (CCA)**

Chosen ciphertext security

- Adversary's power: **CPA** and **CCA**
 - Can encrypt any message of her choice
 - Can decrypt any message of her choice *other than some challenge*
 - (still conservative modeling of real life)
- Adversary's goal: **break semantic security**
 - Learn about the PT from the CT

Chosen ciphertext security (def)

- Let $\zeta = (E, D)$ be a cipher defined over (K, M, C)
- For $b \in \{0, 1\}$ define experiments $\text{EXP}(b)$ as




Chosen ciphertext security (def)

- Def. Cipher $\xi = (E, D)$ is CCA secure if for all efficient adversaries A $\text{Adv}_{\text{CCA}}[A, \xi]$ is negligible.
$$\text{Adv}_{\text{CCA}}[A, \xi] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$
- Thm. A cipher that provides AE is also CCA secure.
- Implication. AE provides confidentiality against an active adversary that can decrypt some ciphertexts.
- Limitations
 - AE does not prevent replay attacks
 - Does not account for side channels attacks (timing)


Ex: AES-CTR is not CCA secure

- Recall
 - AES-CTR is effectively a stream cipher
 - Malleability of stream ciphers




Encrypt then MAC

- MAC computed over cipher text
- Used in IPsec, always provides AE
 - Use separate and independent keys




Encrypt and MAC

- MAC computed over plain text and sent unencrypted
- Used in SSH
- Use separate and independent keys




MAC then encrypt

- MAC computed over plain text and then encrypted before sending
- Used in TLS/SSL
- Use separate and independent keys




Three AE approaches


Encrypt then MAC



Encrypt and MAC




MAC then encrypt



AE: Standardized solutions

- Galois/Counter Mode (GCM)
 - CTR mode encryption then CW-MAC
 - Made popular by Intel's PCLMULQDQ instruction
- CBC-MAC then CTR mode encryption (CCM)
- EAX
- All support ***authenticated encryption with associated data (AEAD)***



Authenticated Encryption

Contents

- Ciphertext integrity
- AE definitions
- Chosen Ciphertext Attack
- Constructions
 - Encrypt-then-MAC
 - Encrypt-and-MAC
 - MAC-then-Encrypt

Authenticated Encryption (AE)

- Everything demonstrated so far provides
 - either integrity
 - or confidentiality (security against eavesdropping)
- CPA security does not provide secrecy against active attacks (where an attacker can tamper with ciphertext)
 - ➔ If you require integrity → **MAC**
 - ➔ If you require integrity and confidentiality → **AE**

AE: Desired properties

- An authenticated encryption system $\zeta = (E, D)$ is a cipher where

as usual $E : K \times M \times N \rightarrow C$

but $D : K \times C \times N \rightarrow M \cup \{\perp\}$ $\perp \notin M$


Nonce

CT is invalid
(rejected)

- Security: the system must provide
 - **semantic security under CPA**, and
 - **ciphertext integrity**
 - an adversary cannot create a new valid CT (such that would decrypt properly)

Ciphertext integrity (def)

Let $\zeta = (E, D)$ be a cipher with message space M



$\downarrow b \in \{0, 1\}$

$b = 1$ if $D(k, c) \neq \perp$ and $c \notin \{c_1 \dots c_q\}$

$b = 0$ otherwise

Def: $\zeta = (E, D)$ has **ciphertext integrity** if for all “efficient” adversaries A : $\text{Adv}_{\text{CI}}[A, \zeta]$ is “negligible”.

$$\text{Adv}_{\text{CI}}[A, \zeta] = \Pr[\text{Chal. outputs } 1]$$

Authenticated Encryption

- Def: A cipher $\zeta = (E, D)$ **provides authenticated encryption (AE)** if it is
 - 1) semantically secure under CPA, and
 - 2) has ciphertext integrity.
- Do the following ciphers provide AE:
 - AES-CBC,
 - AES-CTR,
 - RC4?
- Why?

Authenticated Encryption

- Implication 1: Authenticity



- An attacker cannot create a new valid $c \notin \{c_1 \dots c_q\}$
- If message decrypts properly ($D(k, c) \neq \perp$), it must have come from someone who knows secret key k
 - But it could be a replay

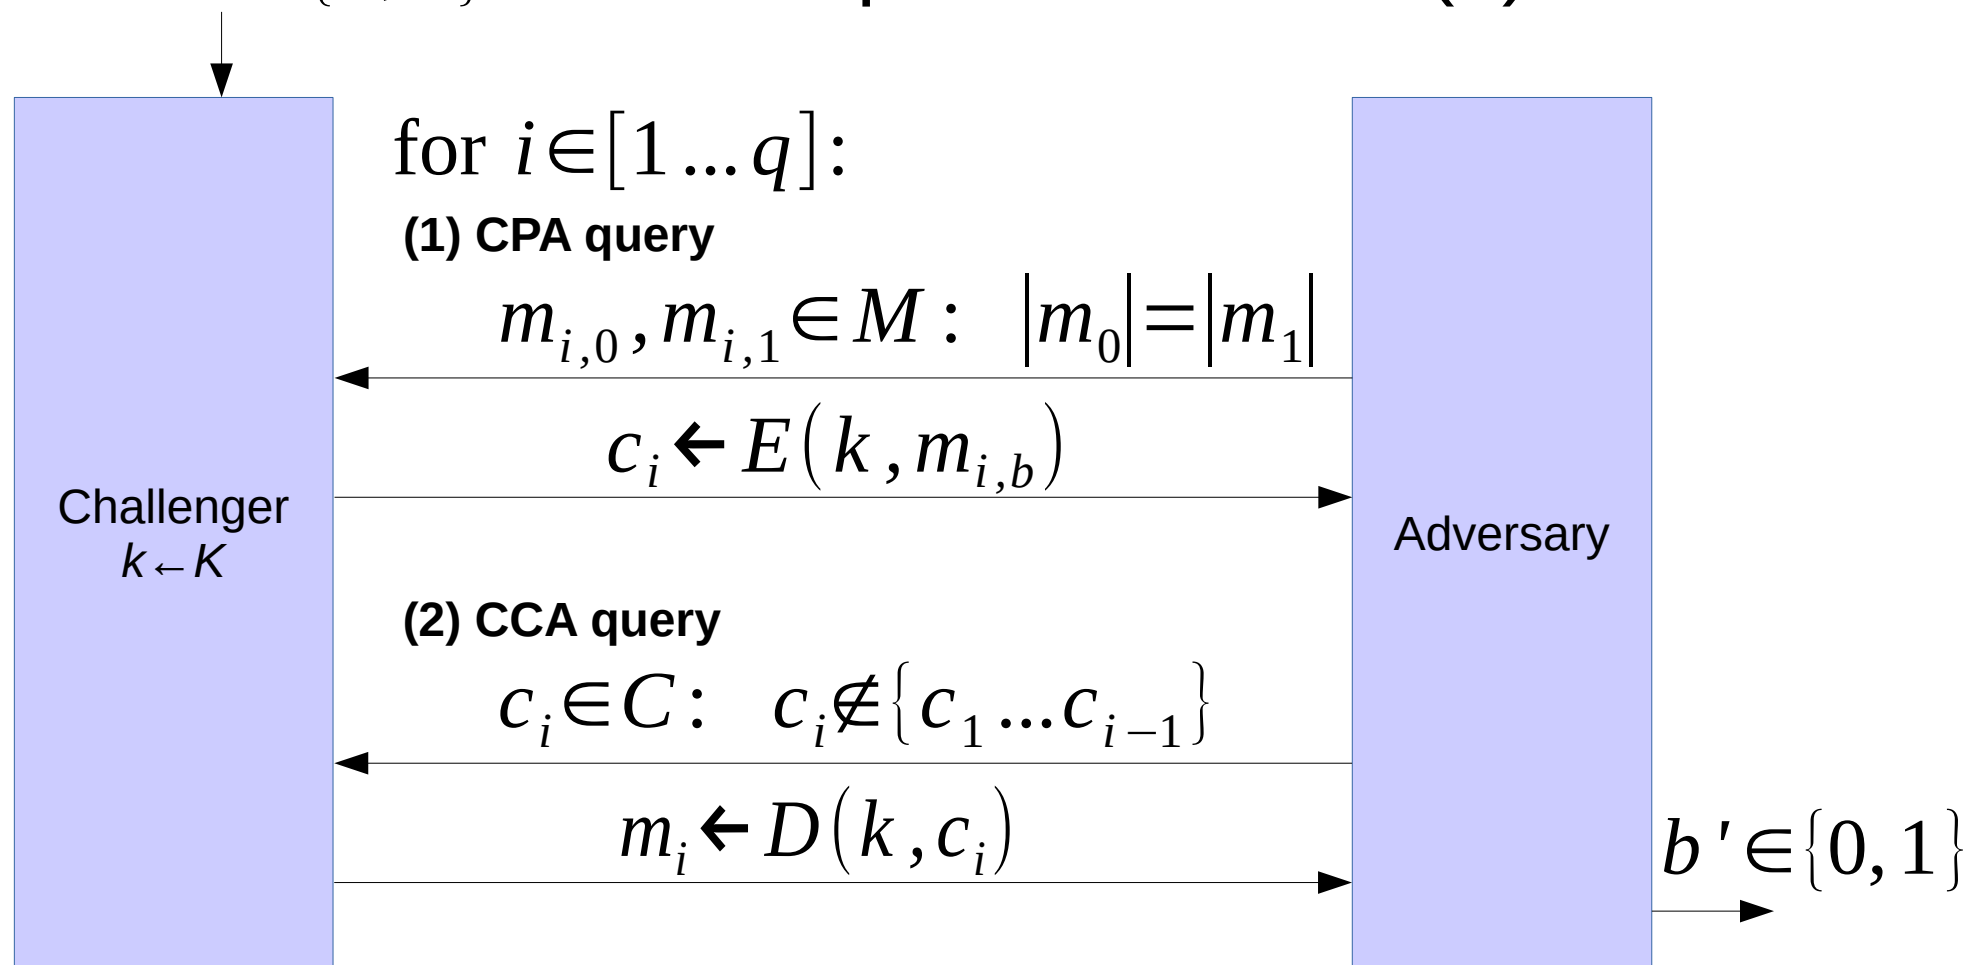
- Implication 2: Security against **chosen ciphertext attack (CCA)**

Chosen ciphertext security

- Adversary's power: **CPA** and **CCA**
 - Can encrypt any message of her choice
 - Can decrypt any message of her choice *other than some challenge*
 - (still conservative modeling of real life)
- Adversary's goal: **break semantic security**
 - Learn about the PT from the CT

Chosen ciphertext security (def)

- Let $\zeta = (E, D)$ be a cipher defined over (K, M, C)
- For $b \in \{0, 1\}$ define experiments $\text{EXP}(b)$ as




Chosen ciphertext security (def)

- Def. Cipher $\xi = (E, D)$ is CCA secure if for all efficient adversaries $A\text{Adv}_{\text{CCA}}[A, \xi]$ is negligible.
$$\text{Adv}_{\text{CCA}}[A, \xi] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$
- Thm. A cipher that provides AE is also CCA secure.
- Implication. AE provides confidentiality against an active adversary that can decrypt some ciphertexts.
- Limitations
 - AE does not prevent replay attacks
 - Does not account for side channels attacks (timing)


Ex: AES-CTR is not CCA secure

- Recall
 - AES-CTR is effectively a stream cipher
 - Malleability of stream ciphers




Encrypt then MAC

- MAC computed over cipher text
- Used in IPsec, always provides AE
 - Use separate and independent keys




Encrypt and MAC

- MAC computed over plain text and sent unencrypted
- Used in SSH
- Use separate and independent keys




MAC then encrypt

- MAC computed over plain text and then encrypted before sending
- Used in TLS/SSL
- Use separate and independent keys




Three AE approaches


Encrypt then MAC



Encrypt and MAC




MAC then encrypt



AE: Standardized solutions

- Galois/Counter Mode (GCM)
 - CTR mode encryption then CW-MAC
 - Made popular by Intel's PCLMULQDQ instruction
- CBC-MAC then CTR mode encryption (CCM)
- EAX
- All support ***authenticated encryption with associated data (AEAD)***




Public key encryption

Index

- Public-key ciphers overview
- Security definitions
 - CPA-security
 - CCA-security
- Trapdoor functions and permutations (TDF, TDP)
 - Encryption schemes from TDF (ISO)
- Example TDP: RSA
 - Definition
 - RSA in practice
 - Security of RSA

Public key encryption

- Each party uses a key pair: $k = (pk, sk)$
- Public key is given to everyone, secret is kept hidden



Public key encryption: usage

- Communication session set-up
 - A process where Alice and Bob agree upon a shared secret
- Non-interactive applications
 - E.g. email
 - Typically, PKs are long-lived, symmetric keys are ephemeral
 - (But the sender needs to know recipient's PK in advance – need PKI)

Public key encryption: def

Def. A public-key encryption system is triple of algs. (G, E, D)

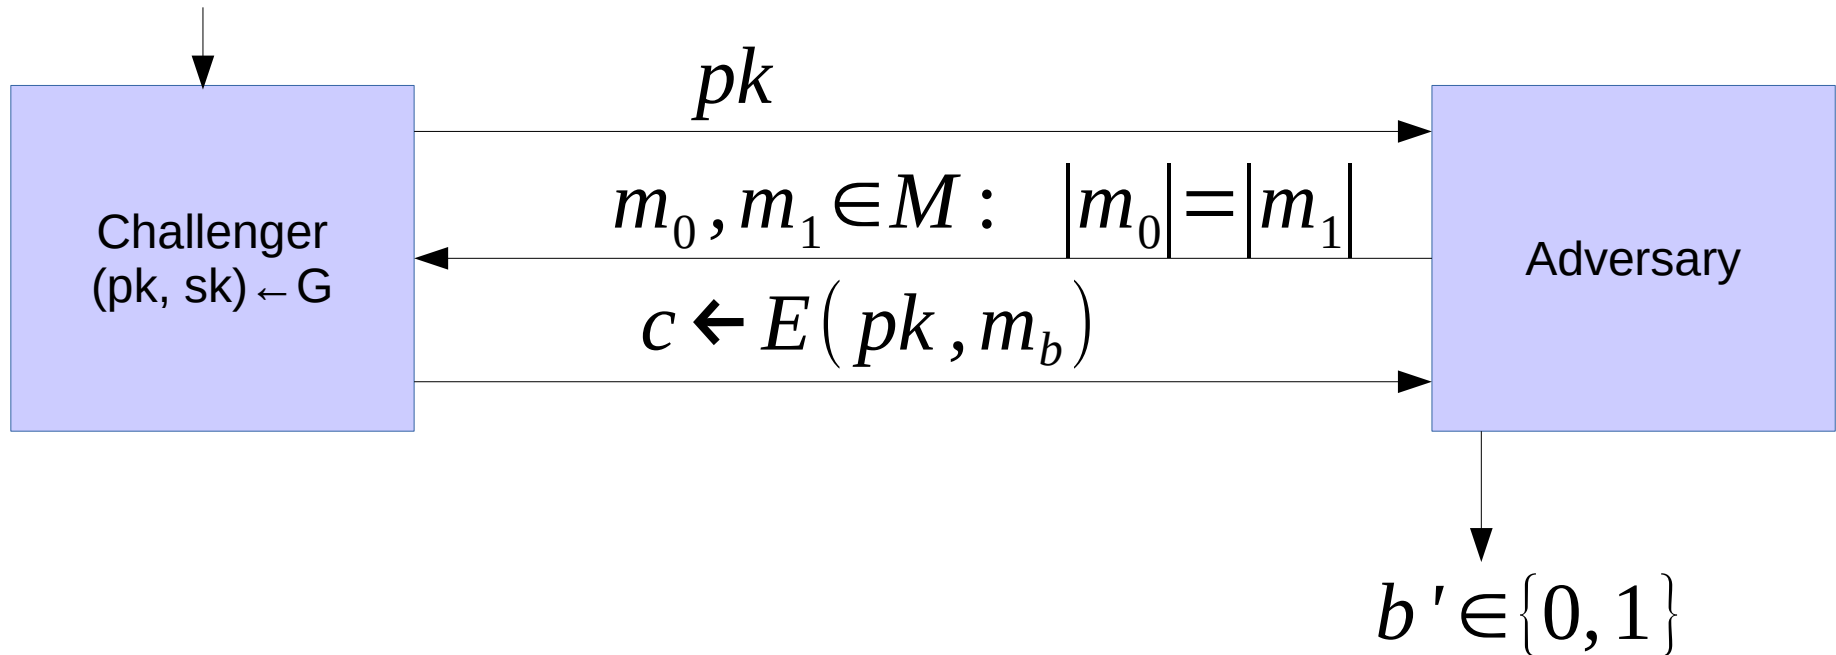
- $G()$ rand. alg. generates key pairs (pk, sk)
- $E(pk, m)$ rand. alg. takes $m \in M$ and returns $c \in C$
- $D(sk, c)$ det. alg. takes $c \in C$ and returns $m \in M$ or \perp

such that $\forall (pk, sk)$ output by G :

$$\forall m \in M : D(sk, E(pk, m)) = m$$

Semantic security (def)

Let $\zeta = (G, E, D)$ be a public key encryption system.
For $b \in \{0, 1\}$ define experiments $\text{EXP}(0)$, $\text{EXP}(1)$



Def: $\zeta = (G, E, D)$ is **semantically secure** (aka IND-CPA) if for all eff. adversaries A : $\text{Adv}_{\text{ss}}[A, \zeta]$ is negligible.

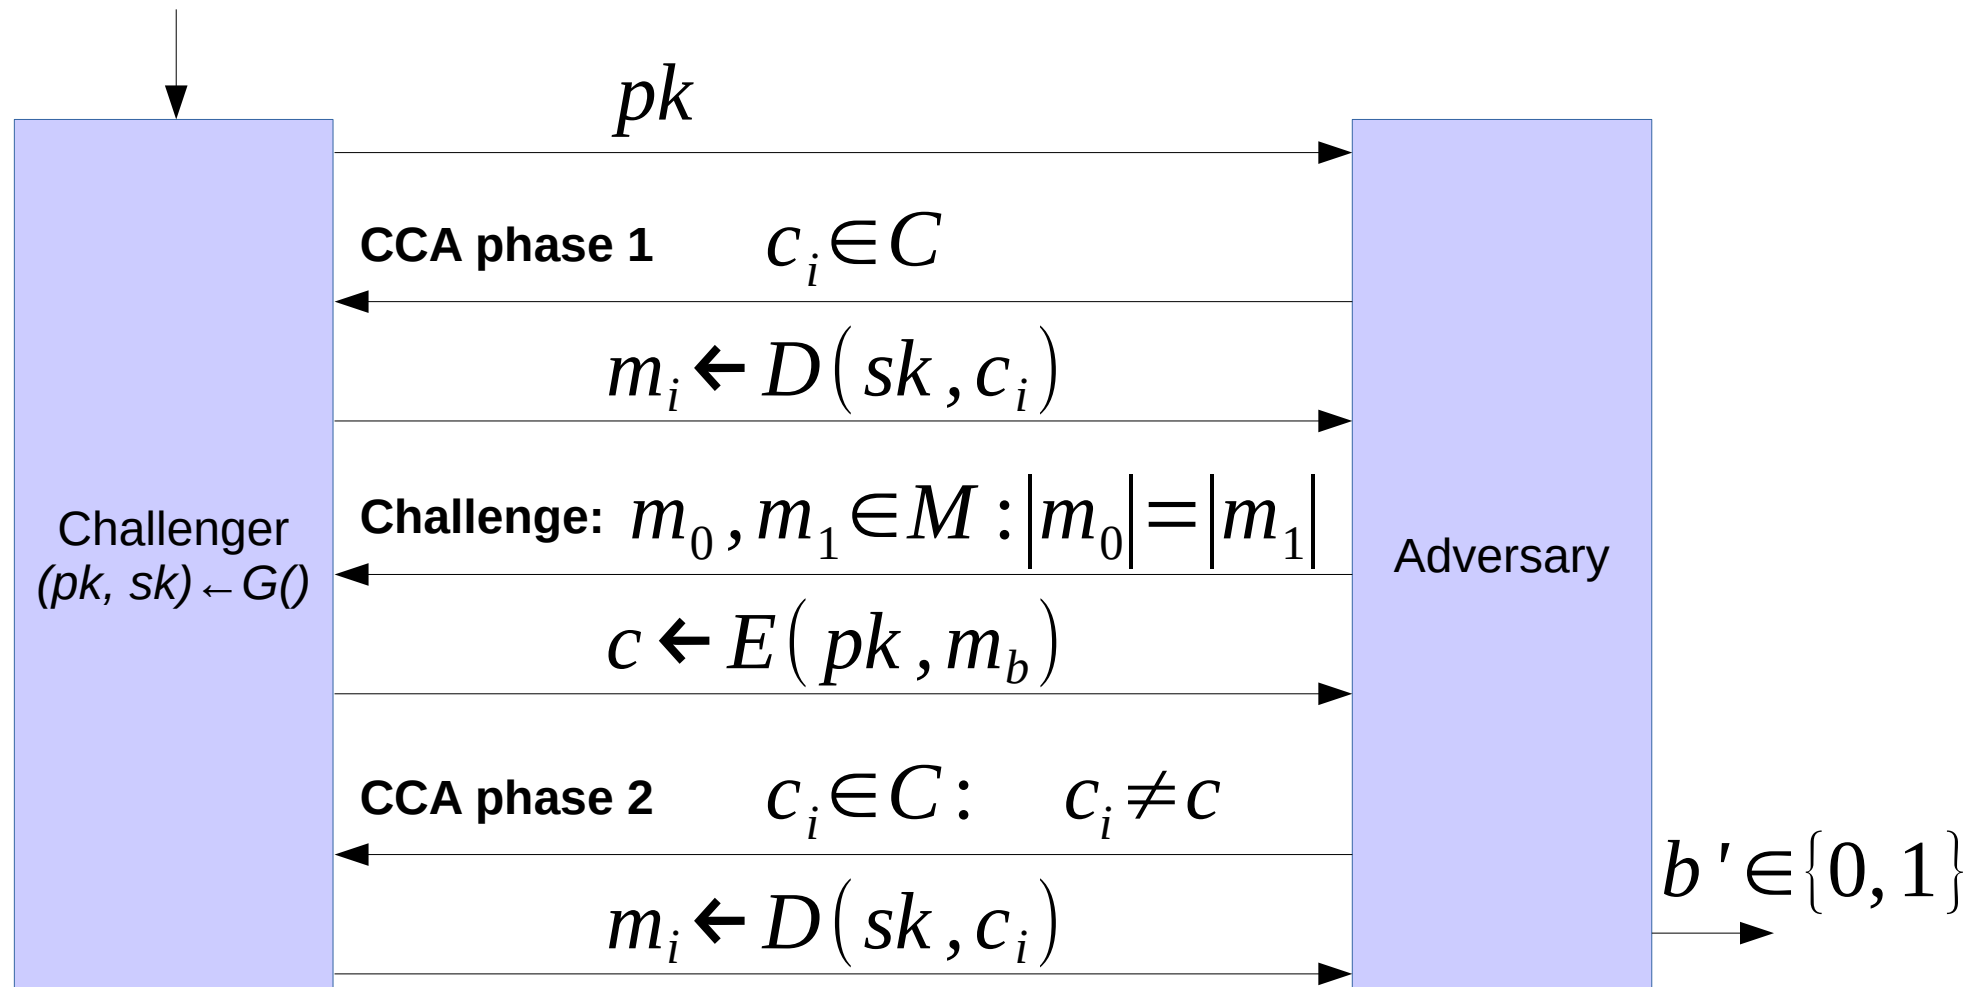
$$\text{Adv}_{\text{ss}}[A, \zeta] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$

Relation to symmetric cipher security

- For symmetric ciphers, we had 2 security definitions
 - One-time security (key used only once) and many-time security (key used many times; CPA)
 - One-time security does not imply many-time security (OTP is broken if used more than once)
- Public key encryption
 - One-time security \rightarrow many-time security (CPA)
 - Because the adversary can encrypt herself (she knows pk)
 - Public key encryption **must be randomized**

(pub-key) Chosen Ciphertext Security (def)

$\zeta = (G, E, D)$ a pub-key enc. over (M, C) . For $b \in \{0, 1\}$ define experiments $\text{EXP}(b)$:



CCA security

- Def. $\zeta = (G, E, D)$ is CCA secure (aka. IND-CCA) if for all efficient adversaries A : $\text{Adv}_{\text{CCA}}[A, \zeta]$ is negligible.

$$\text{Adv}_{\text{CCA}}[A, \zeta] := |\Pr[\text{EXP}(0) = 1] - \Pr[\text{EXP}(1) = 1]|$$

- Recall: A secure symmetric cipher provides AE, when it has CPA security and ciphertext integrity
 - Attacker cannot create new ciphertexts (implies CCA security)
- In pub-key setting
 - Attacker knows $pk \rightarrow$ **can** create new ciphertexts
 - Instead: we directly require CCA security
- Next step: Constructing CCA secure pub-key encryption

Trapdoor function (TDF)


- **Def.** A trapdoor function $X \rightarrow Y$ is a triple of eff. algorithms (G, F, F^{-1})
 - **$G()$** : rand. alg. for creating (pk, sk)
 - **$F(pk, -)$** : det. alg. that defines $X \rightarrow Y$
 - **$F^{-1}(sk, -)$** : det. alg. that defines $Y \rightarrow X$
[inverts $F(pk, -)$]

For every (pk, sk) returned by **G**

$$F^{-1}[sk, F(pk, x)] = x$$

Secure TDFs

- TDF (G, F, F^{-1}) is secure if $F(pk, -)$ is *one-way*
 - It can be evaluated but not inverted without sk



- Def. (G, F, F^{-1}) is a secure TDF if for all eff. algs. A : $\text{Adv}_{\text{OW}}[A, F] := \Pr[x = x']$ is negligible.

Pub-key encryption from TDFs

(ISO 18033-2 standard)

- Building blocks
 - (G, F, F^{-1}) – secure TDF $X \rightarrow Y$
 - (E_s, D_s) – symmetric AE cipher over (K, M, C)
 - $H: X \rightarrow K$ – a hash function
- Pub-key enc. system **(G, E, D)**
 - Key generation **G**: same as **G** in TDF

E(pk, m):

$x \xleftarrow{R} X,$ $y \leftarrow F(pk, x)$
 $k \leftarrow H(x),$ $c \leftarrow E_s(k, m)$

return (y, c)

D(sk, (y, c)):

$x \leftarrow F^{-1}(sk, y)$
 $k \leftarrow H(x),$ $m \leftarrow D_s(k, c)$

return m

Pub-key encryption from TDFs

(ISO 18033-2 standard)

$$F(pk, x)$$

$$E_S(H(x), m)$$

Thm. If $(\mathbf{G}, \mathbf{F}, \mathbf{F}^{-1})$ is a secure TDF, if $(\mathbf{E}_S, \mathbf{D}_S)$ provides AE, and if $\mathbf{H}: \mathbf{X} \rightarrow \mathbf{K}$ is a “random oracle”, then $(\mathbf{G}, \mathbf{E}, \mathbf{D})$ is CCA^{ro} secure.

An incorrect use of TDF:

$$\mathbf{E}(pk, m) := \mathbf{F}(pk, m)$$

$$\mathbf{D}(sk, c) := \mathbf{F}^{-1}(sk, c)$$

Such construction results in a deterministic encryption scheme: cannot be semantically secure

Trapdoor permutation (TDP)

- TDP is a triple of eff. algorithms (G, F, F^{-1})
 - $G()$: generates (pk, sk) ; pk defines a function $X \rightarrow X$
 - $F(pk, x)$: evaluates the function at x
 - $F^{-1}(sk, y)$: inverts the function at y using sk

- **Secure TDP**

The function $F(pk, -)$ is one-way without the sk

Arithmetic modulo composites

Let $N = p \cdot q$ where p, q are primes

$$\mathbb{Z}_N = \{0, 1, \dots, N - 1\}$$

$$\mathbb{Z}_N^* = \{\text{invertible elements in } \mathbb{Z}_N\}$$

Facts $x \in \mathbb{Z}_N$ is invertible $\iff \gcd(x, N) = 1$

$$|\mathbb{Z}_N^*| = \varphi(N) = (p - 1)(q - 1) = N - p - q + 1$$

Euler's theorem

$$\forall x \in \mathbb{Z}_N^* : x^{\varphi(N)} = 1 \pmod{N}$$

RSA trapdoor permutation

- $G()$:
 - Choose random primes p, q (~ 1024 bits); $N = p \cdot q$
 - Choose integers e, d such that $e \cdot d = 1 \pmod{\varphi(N)}$
 - Return $pk = (N, e)$, $sk = (N, d)$
- $F(pk, x)$: $\mathbb{Z}_N^* \rightarrow \mathbb{Z}_N^* : \text{RSA}(x) = x^e \pmod{N}$
- $F^{-1}(sk, y)$:
$$\begin{aligned} y^d &= \text{RSA}(x)^d && \pmod{N} \\ &= x^{ed} && \pmod{N} \\ &= x^{k \cdot \varphi(N) + 1} && \pmod{N} \\ &= (x^{\varphi(N)})^k \cdot x && \pmod{N} \\ &= x \end{aligned}$$

RSA trapdoor permutation

RSA assumption: RSA is one-way permutation

For all eff. algs. A :

$$\Pr[A(N, e, y) = \sqrt[e]{y}] < \text{negligible}$$

$p, q \leftarrow n\text{-bit primes}$

$$N = p \cdot q$$

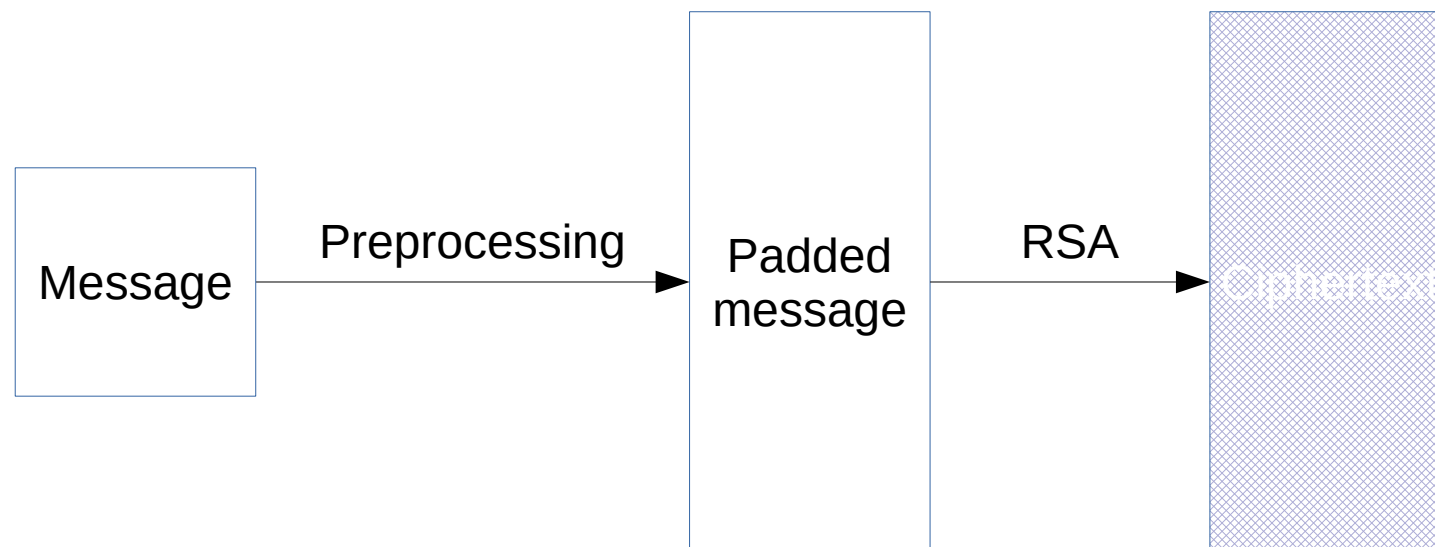
$$y \xleftarrow{R} \mathbb{Z}_N^*$$

Insecure “textbook” RSA


- Encrypting directly with RSA (“textbook” RSA) is insecure
 - $E((N, e), x) := x^e \bmod N$
 - $D((N, d), y) := y^d \bmod N$
- Problem 1: Ciphertext is **malleable**
 - Given ciphertext $c = E((N, e), m)$ an attacker can create $c' = c \cdot 2^e \bmod N$
 - The modified ciphertext c' decrypts to $2m \bmod N$
- Problem 2: Encryption is **deterministic**

RSA in practice

- RSA in practice (ISO standard rarely used)
 - Expand the message to the RSA modulus size and add random bits
 - Apply the RSA function




RSA in practice: PKCS1 v1.5



- Resulting value is RSA encrypted
- Widely deployed (HTTPS)
- Attack due to Bleichenbacher (1998)
 - During decryption, the system will signal an error if the decrypted plaintext does not start with 02
 - Enough to completely decrypt the ciphertext
- Solution in RFC 5246
 - set decrypted PT to a random value and *fail later on*
- Generally PKCS1 v1.5 padding should be avoided

RSA in practice: PKCS1: v2.0 (OAEP)

- New preprocessing function: **Optimal asymmetric encryption padding (OAEP)**
- Check pad on decryption
 - Reject CT if invalid
- **Thm.** If RSA is a TDP, then RSA-OAEP is CCA secure if H, G are *random oracles*.
 - In practice we use SHA-256 for H and G



RSA security (informally)

- To invert RSA one-way function, the attacker must extract x from $c = x^e \bmod N$
- How difficult is to compute e 'th root modulo N ?
Currently best known algorithm
 - Step 1: Factor N [difficult]
 - Step 2: Compute e 'th roots modulo p and q [easy]
- Shor's algorithm: a quantum algorithm for integer factorization in polynomial time
 - Unknown if quantum computers can be built

RSA security (informally)

- Security of public key system should be comparable to security of symmetric cipher

Cipher key size	RSA modulus size [in modulo primes]
80	1024
128	3072
256	15360


Key Exchange

Contents

- Key management problem
- On-line Trusted Third Parties
- The Diffie-Hellman protocol
- Public key cryptography
- Digital signatures
- Key derivation
- Final words


Key management

- Storing mutual secret keys is difficult
- In a universe of n users, each user requires $O(n)$ keys




On-line Trusted Third Party (TTP)

- Every user has to manage only **a single key**
 - The one used to communicate with TTP
- Upon request, the TTP generates shared secret keys for user sessions



TTP: Generating keys (toy protocol)

- Bob wants a shared secret with Alice



TTP: Security

- An eavesdropper sees
 - $m_1 = E(k_B, 'B, A' || k_{AB})$
 - $m_2 = E(k_A, 'B, A' || k_{AB})$
- Since (E, D) is CCA secure, she learns nothing about k_{AB}
- Issues
 - TTP needed for all key exchanges
 - TTP knows all user and all session keys
 - Replay attacks possible
- Basis of Kerberos

The main issue


- Can we generate shared keys without an **on-line** TTP?
 - YES!
- Entrance of public-key cryptography
- Two most widely known constructions
 - Diffie-Hellman protocol (1976)
 - RSA crypto system (1977)

Diffie-Hellman protocol


- Stems from hard problems in algebra
- Alice and Bob want to establish a shared secret in the presence of an eavesdropper
- Security against eavesdropping only




Alice



+




=




Common paint

Secret colours


Bob



+




=




Public transport

(assume
that mixture separation
is expensive)




+




=




Common secret



+



=



Diffie-Hellman protocol (informally)

- Fix a large prime **p** (600 digits ~ 2kbits long)
- Fix an integer **g** in **G** = {1 ... p-1} such that **g** is a primitive root modulo p (generator)
 - Raising **g** to powers of 0 to p-2 generates all values in {1 ... p-1}

Picks a random **a**
in {1 ... p-1}



ALICE

Picks a random **b**
in {1 ... p-1}



BOB

$$A = g^a \mod p$$

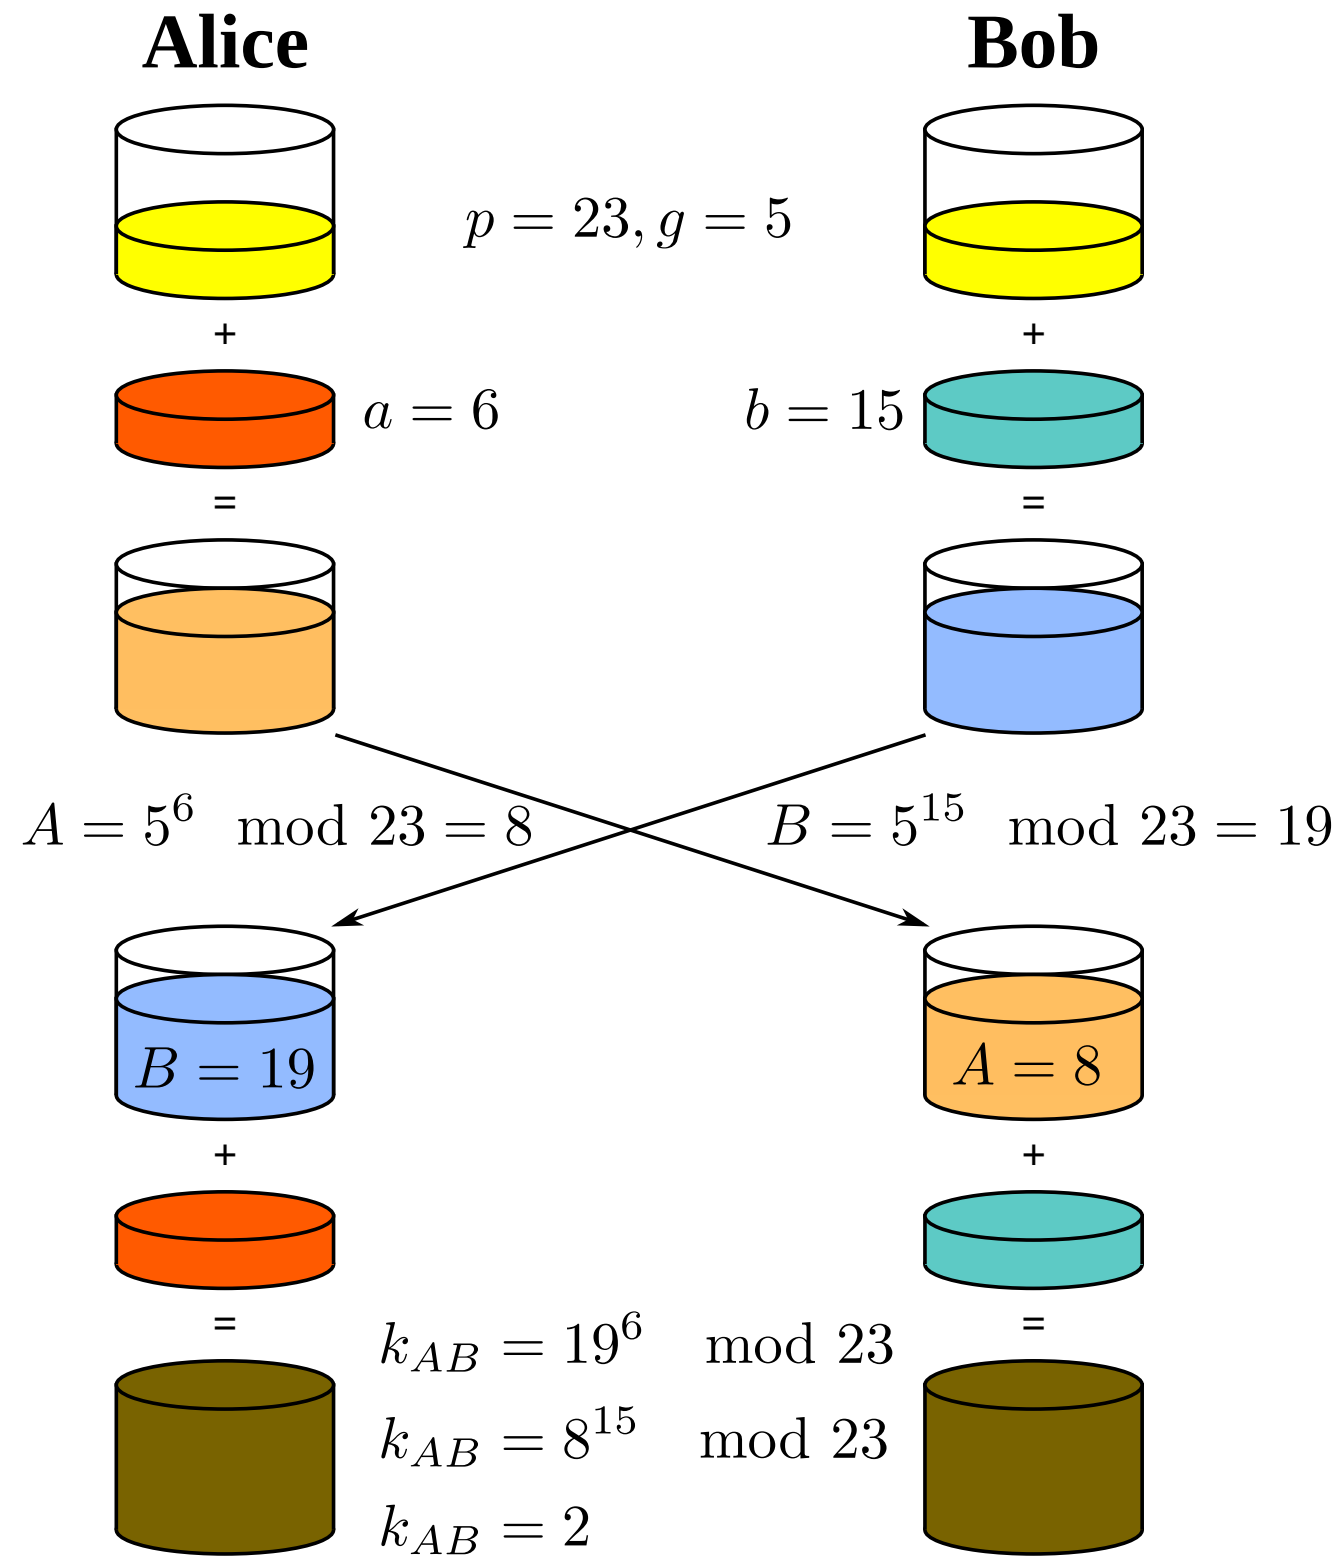
$$B = g^b \mod p$$

$$k_{AB} = g^{ab} \mod p$$

$$B^a = (g^b)^a = g^{ab} \mod p$$

$$A^b = (g^a)^b = g^{ab} \mod p$$

$5^0 = 1$	$\text{mod } 23$
$5^1 = 5$	$\text{mod } 23$
$5^2 = 2$	$\text{mod } 23$
$5^3 = 10$	$\text{mod } 23$
$5^4 = 4$	$\text{mod } 23$
$5^5 = 20$	$\text{mod } 23$
$5^6 = 8$	$\text{mod } 23$
$5^7 = 17$	$\text{mod } 23$
$5^8 = 16$	$\text{mod } 23$
$5^9 = 11$	$\text{mod } 23$
$5^{10} = 9$	$\text{mod } 23$
$5^{11} = 22$	$\text{mod } 23$
$5^{12} = 18$	$\text{mod } 23$
$5^{13} = 21$	$\text{mod } 23$
$5^{14} = 13$	$\text{mod } 23$
$5^{15} = 19$	$\text{mod } 23$
$5^{16} = 3$	$\text{mod } 23$
$5^{17} = 15$	$\text{mod } 23$
$5^{18} = 6$	$\text{mod } 23$
$5^{19} = 7$	$\text{mod } 23$
$5^{20} = 12$	$\text{mod } 23$
$5^{21} = 14$	$\text{mod } 23$
$5^{22} = 5^0 = 1$	$\text{mod } 23$



Security (informally)

- An eavesdropper sees
 - $p, g, A = g^a \pmod{p}, B = g^b \pmod{p}$
- Can she derive $g^{ab} \pmod{p}$ herself?
- In general, let's define
$$\text{DH}_g(g^a, g^b) = g^{ab} \pmod{p}$$
- How difficult is to compute DH function \pmod{p} ?

Security (informally)

- Suppose p is n bits long
- Best known algorithm (GNFS) computes function DH in $e^{O(\sqrt[3]{n})}$
- How difficult is to break DH compared to breaking a symmetric cipher?

Cipher key size	DH modulus size [in modulo primes]	DH modulus size [Elliptic Curve]
80	1024	160
128	3072	256
256	15360	512

- Slow transition from (mod p) to elliptic curves

DH: Open issues

- Remember: security against eavesdropping only
- An active attacker can break the protocol with the man-in-the-middle attack
 - Reason: exchanges are not authenticated


Public key encryption for key exchange

- Alice and Bob want to establish a shared secret in the presence of an eavesdropper
- Security against eavesdropping only

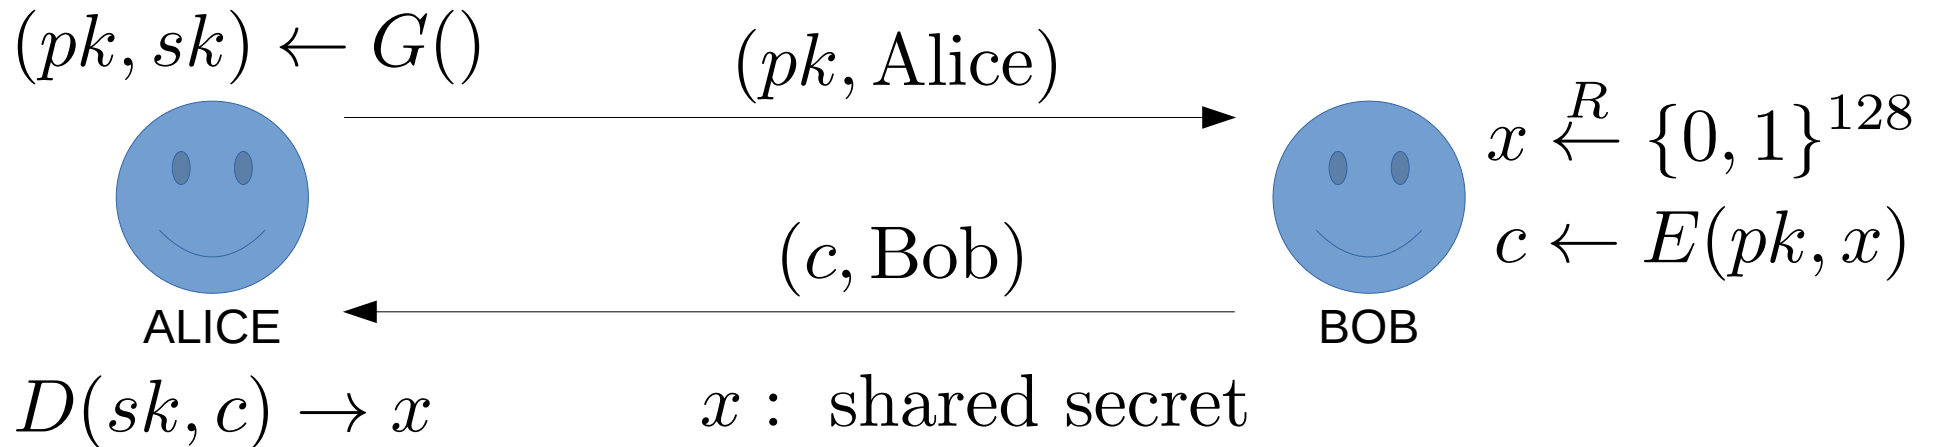


Public key encryption

- Each party uses a key pair: $k = (pk, sk)$
- Public key is given to everyone, secret is kept hidden



Establishing a shared secret



- Adversary sees $pk, E(pk, x)$
- Adversary wants x
- If $\zeta = (G, E, D)$ is sem. secure, the adv. obtains no information about x
- **Security against eavesdropping only:** protocol still vulnerable to man-in-the-middle

Digital signatures

- Preserving integrity in public-key cryptography
 - “MACs” of public-key cryptography
- Idea: The signer signs a message with her secret key.
Anyone can verify the signature using the corresponding public key and thus know:
 - That the message has not been tampered with
 - That the signer indeed signed the message
- Similar to MACs, but digital signatures are
 - **Publicly verifiable**: anyone (with PK) can verify the signature
 - **Non-repudiative**: the signer cannot later deny having signed a particular message

Signature scheme: def.


- **Def:** A signature scheme (G, S, V) is a triple of eff. algs. defined over (M, Z) where:
 - $G()$ is a rand. alg. that generates key pairs (pk, sk)
 - $S(sk, m)$ is an alg. that signs a message $m \in M$ using secret key sk and produces a signature $z \in Z$
 - $V(pk, m, z)$ is a det. alg. that verifies the signature $z \in Z$ of message $m \in M$ using pk and outputs **1** if the signature verifies, or **0** otherwise
- A signature generated by S must always verify by V :
$$\forall (pk, sk), m \in M : \Pr[V(pk, m, S(sk, m)) = 1] = 1$$

Digital signatures: Threat model

- Attacker's power: **Chosen message attack**
 - For $m_1 \dots m_q$ attacker is given $z_i = S(sk, m_i)$
- Attacker's goal: **Existential forgery**
 - Produce a **new** valid (m, z) s. t.
$$m \notin \{m_1 \dots m_q\}$$

→ An adversary cannot produce a valid signature for a new message

Secure digital signature: def.



$b = 1$ if $V(pk, m, z) = 1$ and $m \notin \{m_1 \dots m_q\}$

$b = 0$ otherwise

A signature scheme (G, S, V) is **secure** if for all “efficient” adversaries A :

$$\text{Adv}_{\text{SIG}}[A, I] = \Pr[\text{Chal. outputs } 1]$$

is “negligible”.

Extending the message space

- **Hash-and-sign paradigm**

- *Constructing a signature scheme for large messages from a signature scheme for small messages (and strengthening security)*

- **Thm.** Let (G, S, V) be a secure signature scheme over (M, Z) and let $H: M' \rightarrow M$ be a collision resistant hash function where $|M'| \gg |M|$. Then (G, S', V') is also secure sig. scheme, where:

$$S'(sk, m) := S(sk, H(m))$$

$$V'(pk, m, z) := V(pk, H(m), z)$$

Signatures from TDP: Full Domain Hash

- Building blocks
 - $(\mathbf{G}, \mathbf{F}, \mathbf{F}^{-1})$ – Secure trapdoor permutation (TDP)
 - $\mathbf{F}: \mathbf{X} \rightarrow \mathbf{X}$
 - $\mathbf{H}: \mathbf{M} \rightarrow \mathbf{X}$ – collision resistant hash function
- Full domain (length) hash (FDH)
 - $\mathbf{G}()$ from TDP
 - $S(sk, m) := \mathbf{F}^{-1}(sk, H(m))$
 - $V(pk, m, z) := \begin{cases} 1 & H(m) = \mathbf{F}(pk, z) \\ 0 & \text{otherwise} \end{cases}$

Signatures from TDP: Full Domain Hash

- **Thm.** Let (G, F, F^{-1}) be a secure TDP $X \rightarrow X$ and let $H: M \rightarrow X$ be a collision resistant hash function. Then signature scheme FDH is secure if H is a *random oracle*.
- FDH produces unique signatures: every message has its own signature

Signatures from TDP: Full Domain Hash

- Hashing is required for security; schemes without hashing are insecure. For instance:

$$S(sk, m) := F^{-1}(sk, m) \quad V(pk, m, z) := F(pk, z) == m$$

- **Zero-message attack:** create an existential forgery by picking a random signature, and creating a “message” from it

$$z \xleftarrow{R} Z, m \leftarrow F(pk, z)$$

- **Multiplicative-property attack** (when using RSA)

- Ask for signatures on two messages m_1, m_2

$$z_1 \leftarrow S(sk, m_1), z_2 \leftarrow S(sk, m_2)$$

- Output existential forgery

$$\begin{aligned} m_3 &\leftarrow m_1 \cdot m_2 \\ z_3 &\leftarrow z_1 \cdot z_2 \end{aligned}$$

Signatures from RSA trapdoor


- **$G()$**
 - Choose random primes p, q (~ 1024 bits); $N = p \cdot q$
 - Choose integers e, d such that $e \cdot d = 1 \pmod{\varphi(N)}$
 - Return $pk = (N, e)$, $sk = (N, d)$
- $S((N, d), m) := H(m)^d \pmod N$
- $V((N, e), m, z) := \begin{cases} 1 & H(m) = z^e \pmod N \\ 0 & \text{otherwise} \end{cases}$
- What about **H** ?

RSA Full Domain Hash

- We require $H: M \rightarrow \mathbb{Z}_N^*$
 - The output length of H depends on N ; could be different for every public key
 - Ideally we want the output length of H to be fixed
- **Thm.** Let $H: M \rightarrow Y$ be a collision resistant hash function where $Y = \{1, \dots, 2^{n-2}\}$ and n is the number of bits used to represent N . Then RSA-FDH is secure sig. scheme if H is a random oracle.
 - The bit-length of digests must be of similar length as is the bit-length of the modulus $|Y| \geq N/4$

PKCS1 v1.5 signatures

- Widely deployed (TLS certificates, S/MIME, ...)



- DI** – digest info encodes the name of the used hash function H (SHA*, MD*, ...)
- The resulting value is then signed by raising it to d in $mod N$ (recall, $sk = (N, d)$)
- Not FDH, but partial domain hash
 - No security proof; also no known substantial attacks
 - Issue with proving: $H(m)$ maps to a small subset of \mathbb{Z}_N^*

Probabilistic Signature Scheme (PSS)


- Randomizes the signature with a public random value s called **salt**
- $S((N, d), m, s) := [H(s||m) || MGF[H(s||m)] \oplus s]^d \bmod N$
 - MGF – mask generating function that extends the hash size to the full modulus size
- $V((N, e), m, z, s) := \begin{cases} 1 & H(s||m) || MGF[H(s||m)] \oplus s = z^e \bmod N \\ 0 & \text{otherwise} \end{cases}$
 - Provably secure in random oracle model
 - Part of PKCS1 v2.1

Digital Signature Standard (DSS)

- NIST (FIPS 186)
 - Also called Digital Signature Algorithm (**DSA**)
- Relies on the hardness of Dlog
- No known proof of security
 - But also no serious attacks found
- Has an equivalent in elliptic curves (**ECDSA**)


Deriving many keys from one

- **Scenario:** we obtain a single source key (SK)
 - From a hardware random number generator
 - From a key exchange protocol
- We need many keys to secure the session
 - Unidirectional keys, MAC/encryption keys
- **Goal:** generate many keys from a single SK
 - KDF – key derivation function



Deriving many keys from one

- Three cases
 - 1) SK is uniform in key space
 - 2) SK is non-uniform in key space
 - 3) SK is a password



Key derivation: (1) SK is uniform

- Let PRF $F: K \times X \rightarrow \{0, 1\}^n$
- If source key is uniform in K :

$$KDF(sk, ctx, l) := F(sk, ctx||0) || F(sk, ctx||1) || \dots || F(sk, ctx||l)$$

- **ctx**: a string unique to every application
 - Assures that two applications derive independent keys even if they sample the same source key

Key derivation: (2) SK is non-uniform

- The KDF can be directly used only when SK is uniform
 - If SK is not uniform, the PRF output may not look random
- Reasons for non-uniformity of SK
 - Hardware RNG may be *biased*
 - Key-exchange protocol may produce a key that is *uniform in some subset* of K

Key derivation: (2) SK is non-uniform

- **Extract-then-Expand paradigm**
 - Step 1) Use an **extractor** and **SK** to extract a pseudo-random key **k** that is uniform in key space
 - Use **salt**: a fixed public (non-secret) random string
 - Step 2) expand **k** with KDF
- **HKDF** – a KDF from HMAC
 - Step 1) $k \leftarrow \text{HMAC}(\text{salt}, \text{SK})$
 - Step 2) Expand as you would with uniform keys, but use HMAC for PRF and **k** for key
 - <https://tools.ietf.org/html/rfc5869>

Key derivation: (3) SK is a password

- Particular care needed when deriving keys from passwords
 - HKDF unsuitable here: passwords have low entropy
 - Derived keys will be vulnerable to dictionary attack
- General idea: add **salt** and **slow down hashing**
- **PBKDF** – password-based KDF
 - PKCS #5 v2.0 and <https://tools.ietf.org/html/rfc2898>
 - Iterate hash function many times



Final words

- Cryptography is a powerful tool, but it is too easy to use it incorrectly
 - Systems work, but could be easily attacked
- To reduce the probability of making mistakes
 - Have others review your design and code
 - Never invent your own primitives (ciphers, MACs, modes of operation, ...)
 - Avoid implementing your own cryptographic operations
 - E.g. instead of combining AES-CTR and HMAC, prefer AES-GCM


Key Exchange

Contents

- Key management problem
- On-line Trusted Third Parties
- The Diffie-Hellman protocol
- Public key cryptography
- Digital signatures
- Key derivation
- Final words


Key management

- Storing mutual secret keys is difficult
- In a universe of n users, each user requires $O(n)$ keys



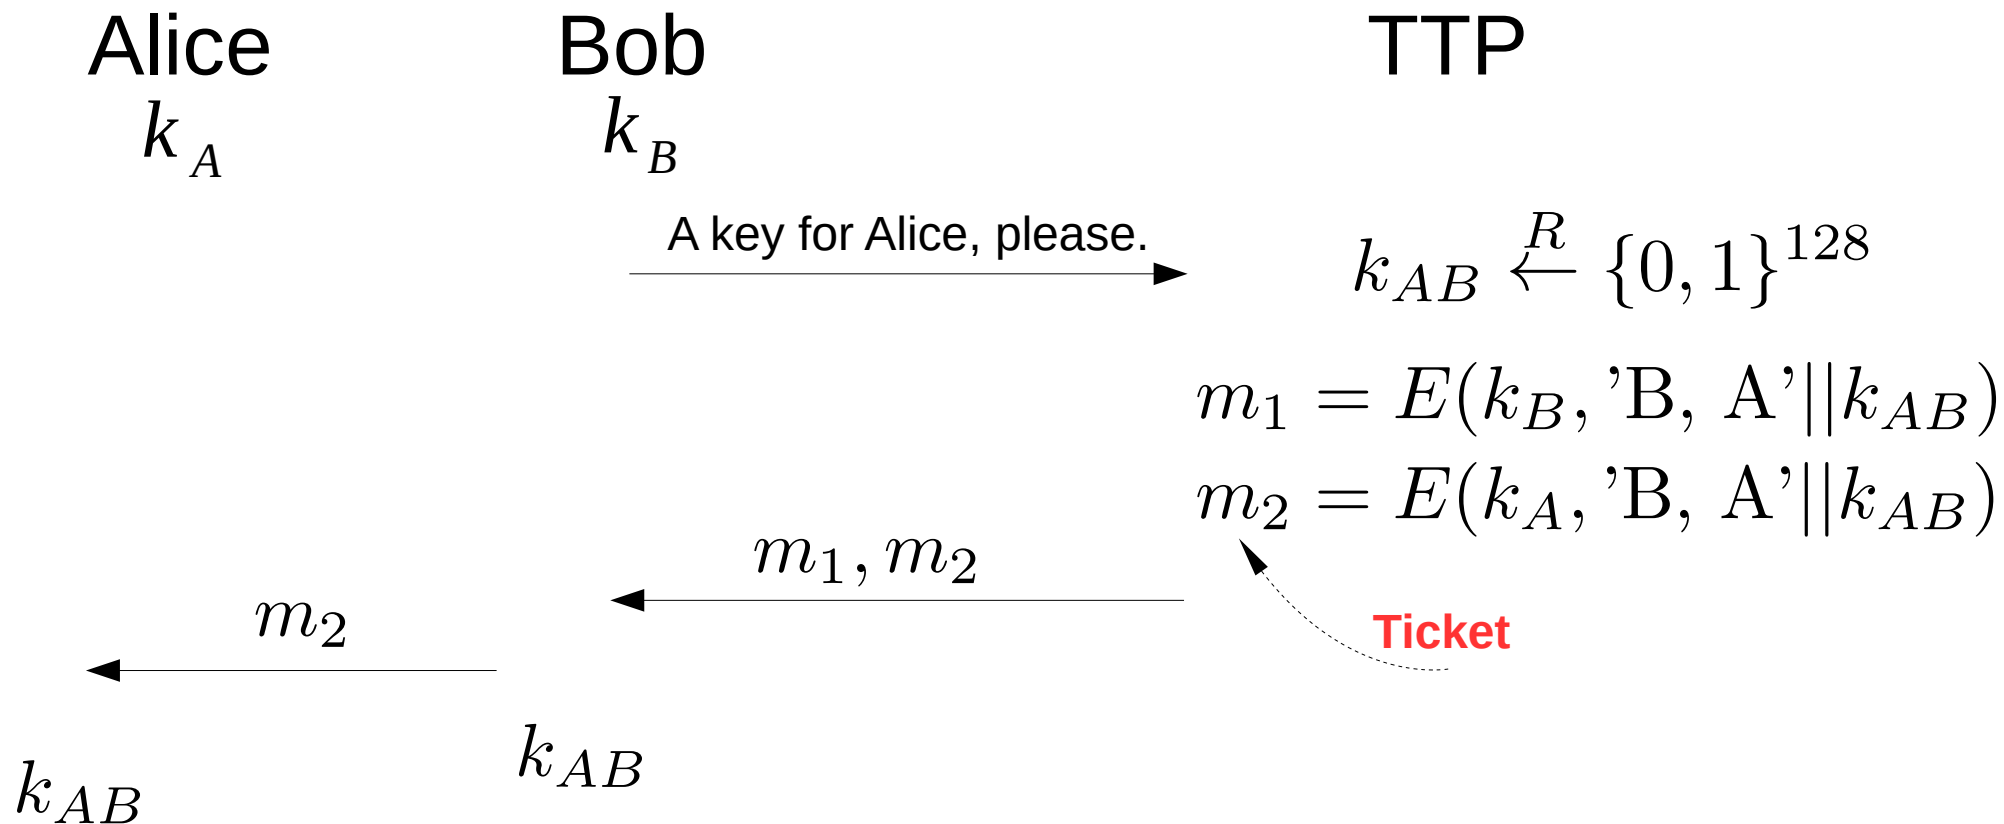
On-line Trusted Third Party (TTP)

- Every user has to manage only **a single key**
 - The one used to communicate with TTP
- Upon request, the TTP generates shared secret keys for user sessions



TTP: Generating keys (toy protocol)

- Bob wants a shared secret with Alice



TTP: Security

- An eavesdropper sees
 - $m_1 = E(k_B, 'B, A' || k_{AB})$
 - $m_2 = E(k_A, 'B, A' || k_{AB})$
- Since (E, D) is CCA secure, she learns nothing about k_{AB}
- Issues
 - TTP needed for all key exchanges
 - TTP knows all user and all session keys
 - Replay attacks possible
- Basis of Kerberos

The main issue


- Can we generate shared keys without an **on-line** TTP?
 - YES!
- Entrance of public-key cryptography
- Two most widely known constructions
 - Diffie-Hellman protocol (1976)
 - RSA crypto system (1977)

Diffie-Hellman protocol


- Stems from hard problems in algebra
- Alice and Bob want to establish a shared secret in the presence of an eavesdropper
- Security against eavesdropping only




Alice



+




=




Common paint

Secret colours


Bob



+




=




Public transport

(assume
that mixture separation
is expensive)




+




=




Common secret



+



=



Diffie-Hellman protocol (informally)


- Fix a large prime p (600 digits \sim 2kbits long)
- Fix an integer g in $G = \{1 \dots p-1\}$ such that g is a primitive root modulo p (generator)
 - Raising g to powers of 0 to $p-2$ generates all values in $\{1 \dots p-1\}$

Picks a random a
in $\{1 \dots p-1\}$



ALICE

Picks a random b
in $\{1 \dots p-1\}$



BOB

$$A = g^a \mod p$$


$$B = g^b \mod p$$

$$k_{AB} = g^{ab} \mod p$$

$$B^a = (g^b)^a = g^{ab} \mod p$$

$$A^b = (g^a)^b = g^{ab} \mod p$$

$5^0 = 1$	$\text{mod } 23$
$5^1 = 5$	$\text{mod } 23$
$5^2 = 2$	$\text{mod } 23$
$5^3 = 10$	$\text{mod } 23$
$5^4 = 4$	$\text{mod } 23$
$5^5 = 20$	$\text{mod } 23$
$5^6 = 8$	$\text{mod } 23$
$5^7 = 17$	$\text{mod } 23$
$5^8 = 16$	$\text{mod } 23$
$5^9 = 11$	$\text{mod } 23$
$5^{10} = 9$	$\text{mod } 23$
$5^{11} = 22$	$\text{mod } 23$
$5^{12} = 18$	$\text{mod } 23$
$5^{13} = 21$	$\text{mod } 23$
$5^{14} = 13$	$\text{mod } 23$
$5^{15} = 19$	$\text{mod } 23$
$5^{16} = 3$	$\text{mod } 23$
$5^{17} = 15$	$\text{mod } 23$
$5^{18} = 6$	$\text{mod } 23$
$5^{19} = 7$	$\text{mod } 23$
$5^{20} = 12$	$\text{mod } 23$
$5^{21} = 14$	$\text{mod } 23$
$5^{22} = 5^0 = 1$	$\text{mod } 23$



Security (informally)

- An eavesdropper sees
 - $p, g, A = g^a \pmod{p}, B = g^b \pmod{p}$
- Can she derive $g^{ab} \pmod{p}$ herself?
- In general, let's define
$$\text{DH}_g(g^a, g^b) = g^{ab} \pmod{p}$$
- How difficult is to compute DH function \pmod{p} ?

Security (informally)

- Suppose p is n bits long
- Best known algorithm (GNFS) computes function DH in $e^{O(\sqrt[3]{n})}$
- How difficult is to break DH compared to breaking a symmetric cipher?

Cipher key size	DH modulus size [in modulo primes]	DH modulus size [Elliptic Curve]
80	1024	160
128	3072	256
256	15360	512

- Slow transition from (mod p) to elliptic curves

DH: Open issues

- Remember: security against eavesdropping only
- An active attacker can break the protocol with the man-in-the-middle attack
 - Reason: exchanges are not authenticated


Public key encryption for key exchange

- Alice and Bob want to establish a shared secret in the presence of an eavesdropper
- Security against eavesdropping only

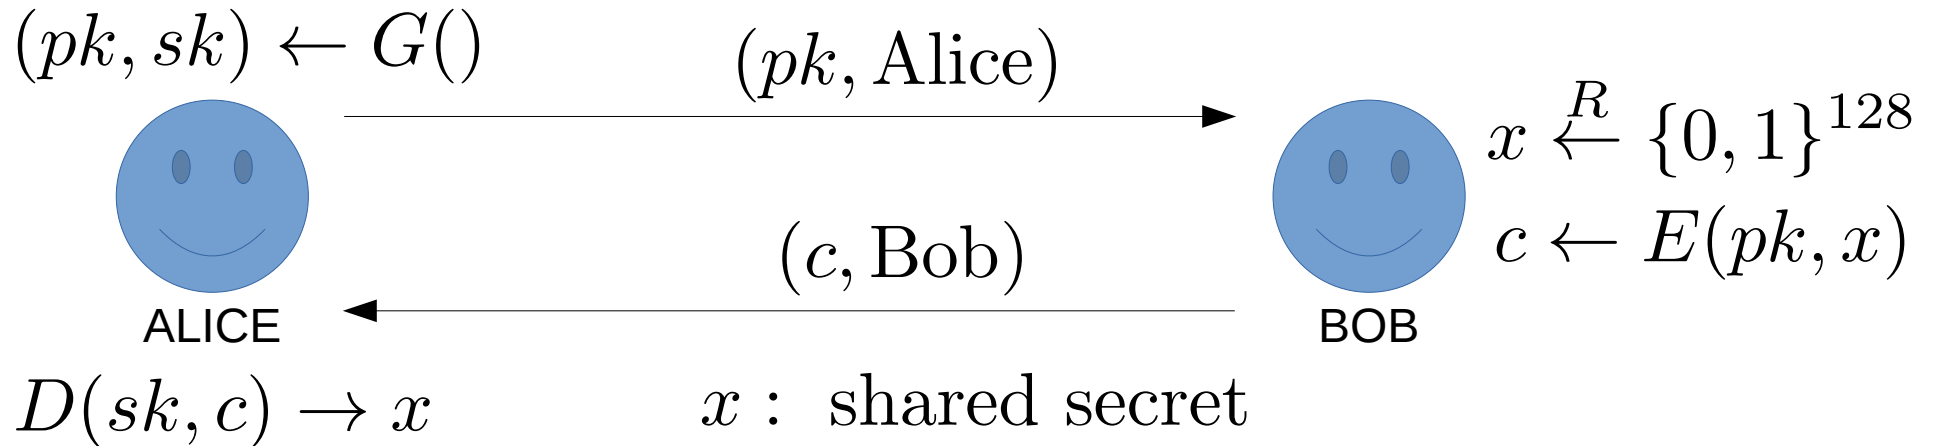


Public key encryption

- Each party uses a key pair: $k = (\text{pk}, \text{sk})$
- Public key is given to everyone, secret is kept hidden



Establishing a shared secret



- Adversary sees $pk, E(pk, x)$
- Adversary wants x
- If $\zeta = (G, E, D)$ is sem. secure, the adv. obtains no information about x
- **Security against eavesdropping only:** protocol still vulnerable to man-in-the-middle

Digital signatures

- Preserving integrity in public-key cryptography
 - “MACs” of public-key cryptography
- Idea: The signer signs a message with her secret key.
Anyone can verify the signature using the corresponding public key and thus know:
 - That the message has not been tampered with
 - That the signer indeed signed the message
- Similar to MACs, but digital signatures are
 - **Publicly verifiable**: anyone (with PK) can verify the signature
 - **Non-repudiative**: the signer cannot later deny having signed a particular message

Signature scheme: def.


- **Def:** A signature scheme (G, S, V) is a triple of eff. algs. defined over (M, Z) where:
 - $G()$ is a rand. alg. that generates key pairs (pk, sk)
 - $S(sk, m)$ is an alg. that signs a message $m \in M$ using secret key sk and produces a signature $z \in Z$
 - $V(pk, m, z)$ is a det. alg. that verifies the signature $z \in Z$ of message $m \in M$ using pk and outputs **1** if the signature verifies, or **0** otherwise
- A signature generated by S must always verify by V :
$$\forall (pk, sk), m \in M: \Pr[V(pk, m, S(sk, m)) = 1] = 1$$

Digital signatures: Threat model

- Attacker's power: **Chosen message attack**
 - For $m_1 \dots m_q$ attacker is given $z_i = S(sk, m_i)$
- Attacker's goal: **Existential forgery**
 - Produce a **new** valid (m, z) s. t.
$$m \notin \{m_1 \dots m_q\}$$

→ An adversary cannot produce a valid signature for a new message

Secure digital signature: def.



$b=1$ if $V(pk, m, z)=1$ and $m \notin \{m_1 \dots m_q\}$

$b=0$ otherwise

A signature scheme (G, S, V) is **secure** if for all “efficient” adversaries A :

$$\text{Adv}_{\text{SIG}}[A, I] = \Pr[\text{Chal. outputs } 1]$$

is “negligible”.

Extending the message space

- **Hash-and-sign paradigm**

- *Constructing a signature scheme for large messages from a signature scheme for small messages (and strengthening security)*

- **Thm.** Let (G, S, V) be a secure signature scheme over (M, Z) and let $H: M' \rightarrow M$ be a collision resistant hash function where $|M'| \gg |M|$. Then (G, S', V') is also secure sig. scheme, where:

$$S'(sk, m) := S(sk, H(m))$$

$$V'(pk, m, z) := V(pk, H(m), z)$$

Signatures from TDP: Full Domain Hash

- Building blocks
 - $(\mathbf{G}, \mathbf{F}, \mathbf{F}^{-1})$ – Secure trapdoor permutation (TDP)
 - $\mathbf{F}: \mathbf{X} \rightarrow \mathbf{X}$
 - $\mathbf{H}: \mathbf{M} \rightarrow \mathbf{X}$ – collision resistant hash function
- Full domain (length) hash (FDH)
 - $\mathbf{G}()$ from TDP
 - $S(sk, m) := F^{-1}(sk, H(m))$
 - $V(pk, m, z) := \begin{cases} 1 & H(m) = F(pk, z) \\ 0 & \text{otherwise} \end{cases}$

Signatures from TDP: Full Domain Hash

- **Thm.** Let (G, F, F^{-1}) be a secure TDP $X \rightarrow X$ and let $H: M \rightarrow X$ be a collision resistant hash function. Then signature scheme FDH is secure if H is a *random oracle*.
- FDH produces unique signatures: every message has its own signature

Signatures from TDP: Full Domain Hash

- Hashing is required for security; schemes without hashing are insecure. For instance:

$$S(sk, m) := F^{-1}(sk, m) \quad V(pk, m, z) := F(pk, z) == m$$

- **Zero-message attack:** create an existential forgery by picking a random signature, and creating a “message” from it

$$z \xleftarrow{R} Z, m \leftarrow F(pk, z)$$

- **Multiplicative-property attack** (when using RSA)

- Ask for signatures on two messages m_1, m_2

$$z_1 \leftarrow S(sk, m_1), z_2 \leftarrow S(sk, m_2)$$

- Output existential forgery

$$\begin{aligned} m_3 &\leftarrow m_1 \cdot m_2 \\ z_3 &\leftarrow z_1 \cdot z_2 \end{aligned}$$

Signatures from RSA trapdoor


- **$G()$**
 - Choose random primes p, q (~ 1024 bits); $N = p \cdot q$
 - Choose integers e, d such that $e \cdot d = 1 \pmod{\varphi(N)}$
 - Return $pk = (N, e)$, $sk = (N, d)$
- $S((N, d), m) := H(m)^d \pmod N$
- $V((N, e), m, z) := \begin{cases} 1 & H(m) = z^e \pmod N \\ 0 & \text{otherwise} \end{cases}$
- What about **H** ?

RSA Full Domain Hash

- We require $H: M \rightarrow \mathbb{Z}_N^*$
 - The output length of H depends on N ; could be different for every public key
 - Ideally we want the output length of H to be fixed
- **Thm.** Let $H: M \rightarrow Y$ be a collision resistant hash function where $Y = \{1, \dots, 2^{n-2}\}$ and n is the number of bits used to represent N . Then RSA-FDH is secure sig. scheme if H is a random oracle.
 - The bit-length of digests must be of similar length as is the bit-length of the modulus $|Y| \geq N/4$

PKCS1 v1.5 signatures

- Widely deployed (TLS certificates, S/MIME, ...)



- DI** – digest info encodes the name of the used hash function H (SHA*, MD*, ...)
- The resulting value is then signed by raising it to d in $mod N$ (recall, $sk = (N, d)$)
- Not FDH, but partial domain hash
 - No security proof; also no known substantial attacks
 - Issue with proving: $H(m)$ maps to a small subset of \mathbb{Z}_N^*

Probabilistic Signature Scheme (PSS)


- Randomizes the signature with a public random value s called **salt**
- $S((N, d), m, s) := [H(s||m) || MGF[H(s||m)] \oplus s]^d \bmod N$
 - MGF – mask generating function that extends the hash size to the full modulus size
- $V((N, e), m, z, s) := \begin{cases} 1 & H(s||m) || MGF[H(s||m)] \oplus s = z^e \bmod N \\ 0 & \text{otherwise} \end{cases}$
 - Provably secure in random oracle model
 - Part of PKCS1 v2.1

Digital Signature Standard (DSS)

- NIST (FIPS 186)
 - Also called Digital Signature Algorithm (**DSA**)
- Relies on the hardness of Dlog
- No known proof of security
 - But also no serious attacks found
- Has an equivalent in elliptic curves (**ECDSA**)


Deriving many keys from one

- **Scenario:** we obtain a single source key (SK)
 - From a hardware random number generator
 - From a key exchange protocol
- We need many keys to secure the session
 - Unidirectional keys, MAC/encryption keys
- **Goal:** generate many keys from a single SK
 - KDF – key derivation function



Deriving many keys from one

- Three cases
 - 1) SK is uniform in key space
 - 2) SK is non-uniform in key space
 - 3) SK is a password



Key derivation: (1) SK is uniform

- Let PRF $F: K \times X \rightarrow \{0, 1\}^n$
- If source key is uniform in K :

$$KDF(sk, ctx, l) := F(sk, ctx||0) || F(sk, ctx||1) || \dots || F(sk, ctx||l)$$

- **ctx**: a string unique to every application
 - Assures that two applications derive independent keys even if they sample the same source key

Key derivation: (2) SK is non-uniform

- The KDF can be directly used only when SK is uniform
 - If SK is not uniform, the PRF output may not look random
- Reasons for non-uniformity of SK
 - Hardware RNG may be *biased*
 - Key-exchange protocol may produce a key that is *uniform in some subset* of K

Key derivation: (2) SK is non-uniform

- **Extract-then-Expand paradigm**
 - Step 1) Use an **extractor** and **SK** to extract a pseudo-random key **k** that is uniform in key space
 - Use **salt**: a fixed public (non-secret) random string
 - Step 2) expand **k** with KDF
- **HKDF** – a KDF from HMAC
 - Step 1) $k \leftarrow \text{HMAC}(\text{salt}, \text{SK})$
 - Step 2) Expand as you would with uniform keys, but use HMAC for PRF and **k** for key
 - <https://tools.ietf.org/html/rfc5869>

Key derivation: (3) SK is a password

- Particular care needed when deriving keys from passwords
 - HKDF unsuitable here: passwords have low entropy
 - Derived keys will be vulnerable to dictionary attack
- General idea: add **salt** and **slow down hashing**
- **PBKDF** – password-based KDF
 - PKCS #5 v2.0 and <https://tools.ietf.org/html/rfc2898>
 - Iterate hash function many times



Final words

- Cryptography is a powerful tool, but it is too easy to use it incorrectly
 - Systems work, but could be easily attacked
- To reduce the probability of making mistakes
 - Have others review your design and code
 - Never invent your own primitives (ciphers, MACs, modes of operation, ...)
 - Avoid implementing your own cryptographic operations
 - E.g. instead of combining AES-CTR and HMAC, prefer AES-GCM